



ENABLING PYTHON USER DEFINED FUNCTIONS (UDFS) IN ACCELERATED APPLICATIONS WITH NUMBA

GRAHAM MARKALL, BRANDON MILLER, MICHAEL YH WANG - NVIDIA RAPIDS

SPEAKER INTRODUCTION - GRAHAM MARKALL

- Software Engineer in RAPIDS at NVIDIA:
 - Numba CUDA target maintainer
 - Supporting cuDF / RAPIDS use cases
- Background in compilers / numerical methods / HPC:
 - GCC, Binutils, GDB, LLVM, ...
 - PDEs, Finite elements, sparse linear solvers,
 - Domain-specific languages for High Performance Computing
- Aim: Make powerful software tools accessible to many people

PROBLEM STATEMENT - THE “IMPEDANCE MISMATCH”

- **Applications:**

- Many CUDA-accelerated applications written in C++



- **Users / developers:**

- Many application users / developers prefer to work in Python



- **How do we:**

- Enable Python users to extend accelerated applications...
- ... whilst retaining *developer productivity*...
- ... and *application performance*?



- **Solution:**

- Use Numba to compile Python user code for CUDA
- *User-Defined Functions (UDFs)*

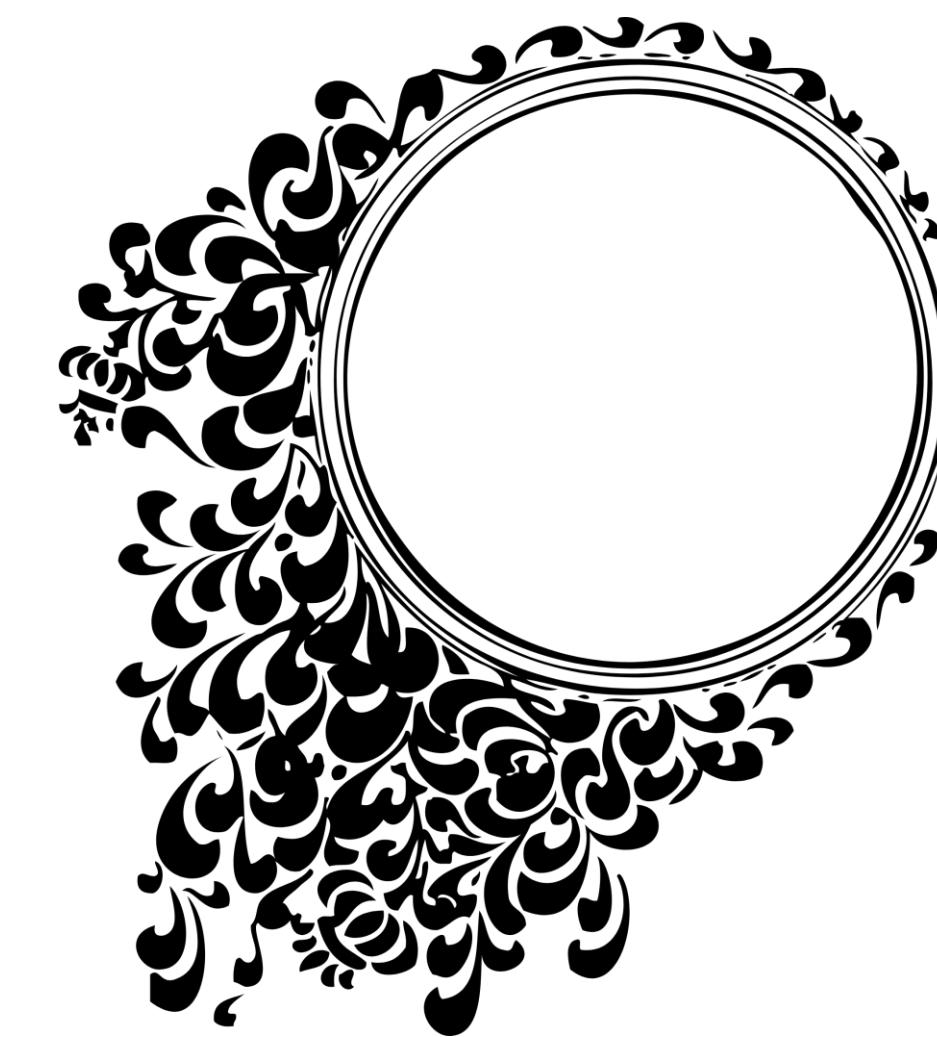
TALK ROADMAP

- Examples:

- cuDF
- PyOptiX
- Filigree

RAPIDS

NVIDIA[®]
OPTIX™



- Numba:

- Overview
- Pipeline & extensions


Numba

- Worked example:

- Implementation
- Outcomes (development ease & performance)

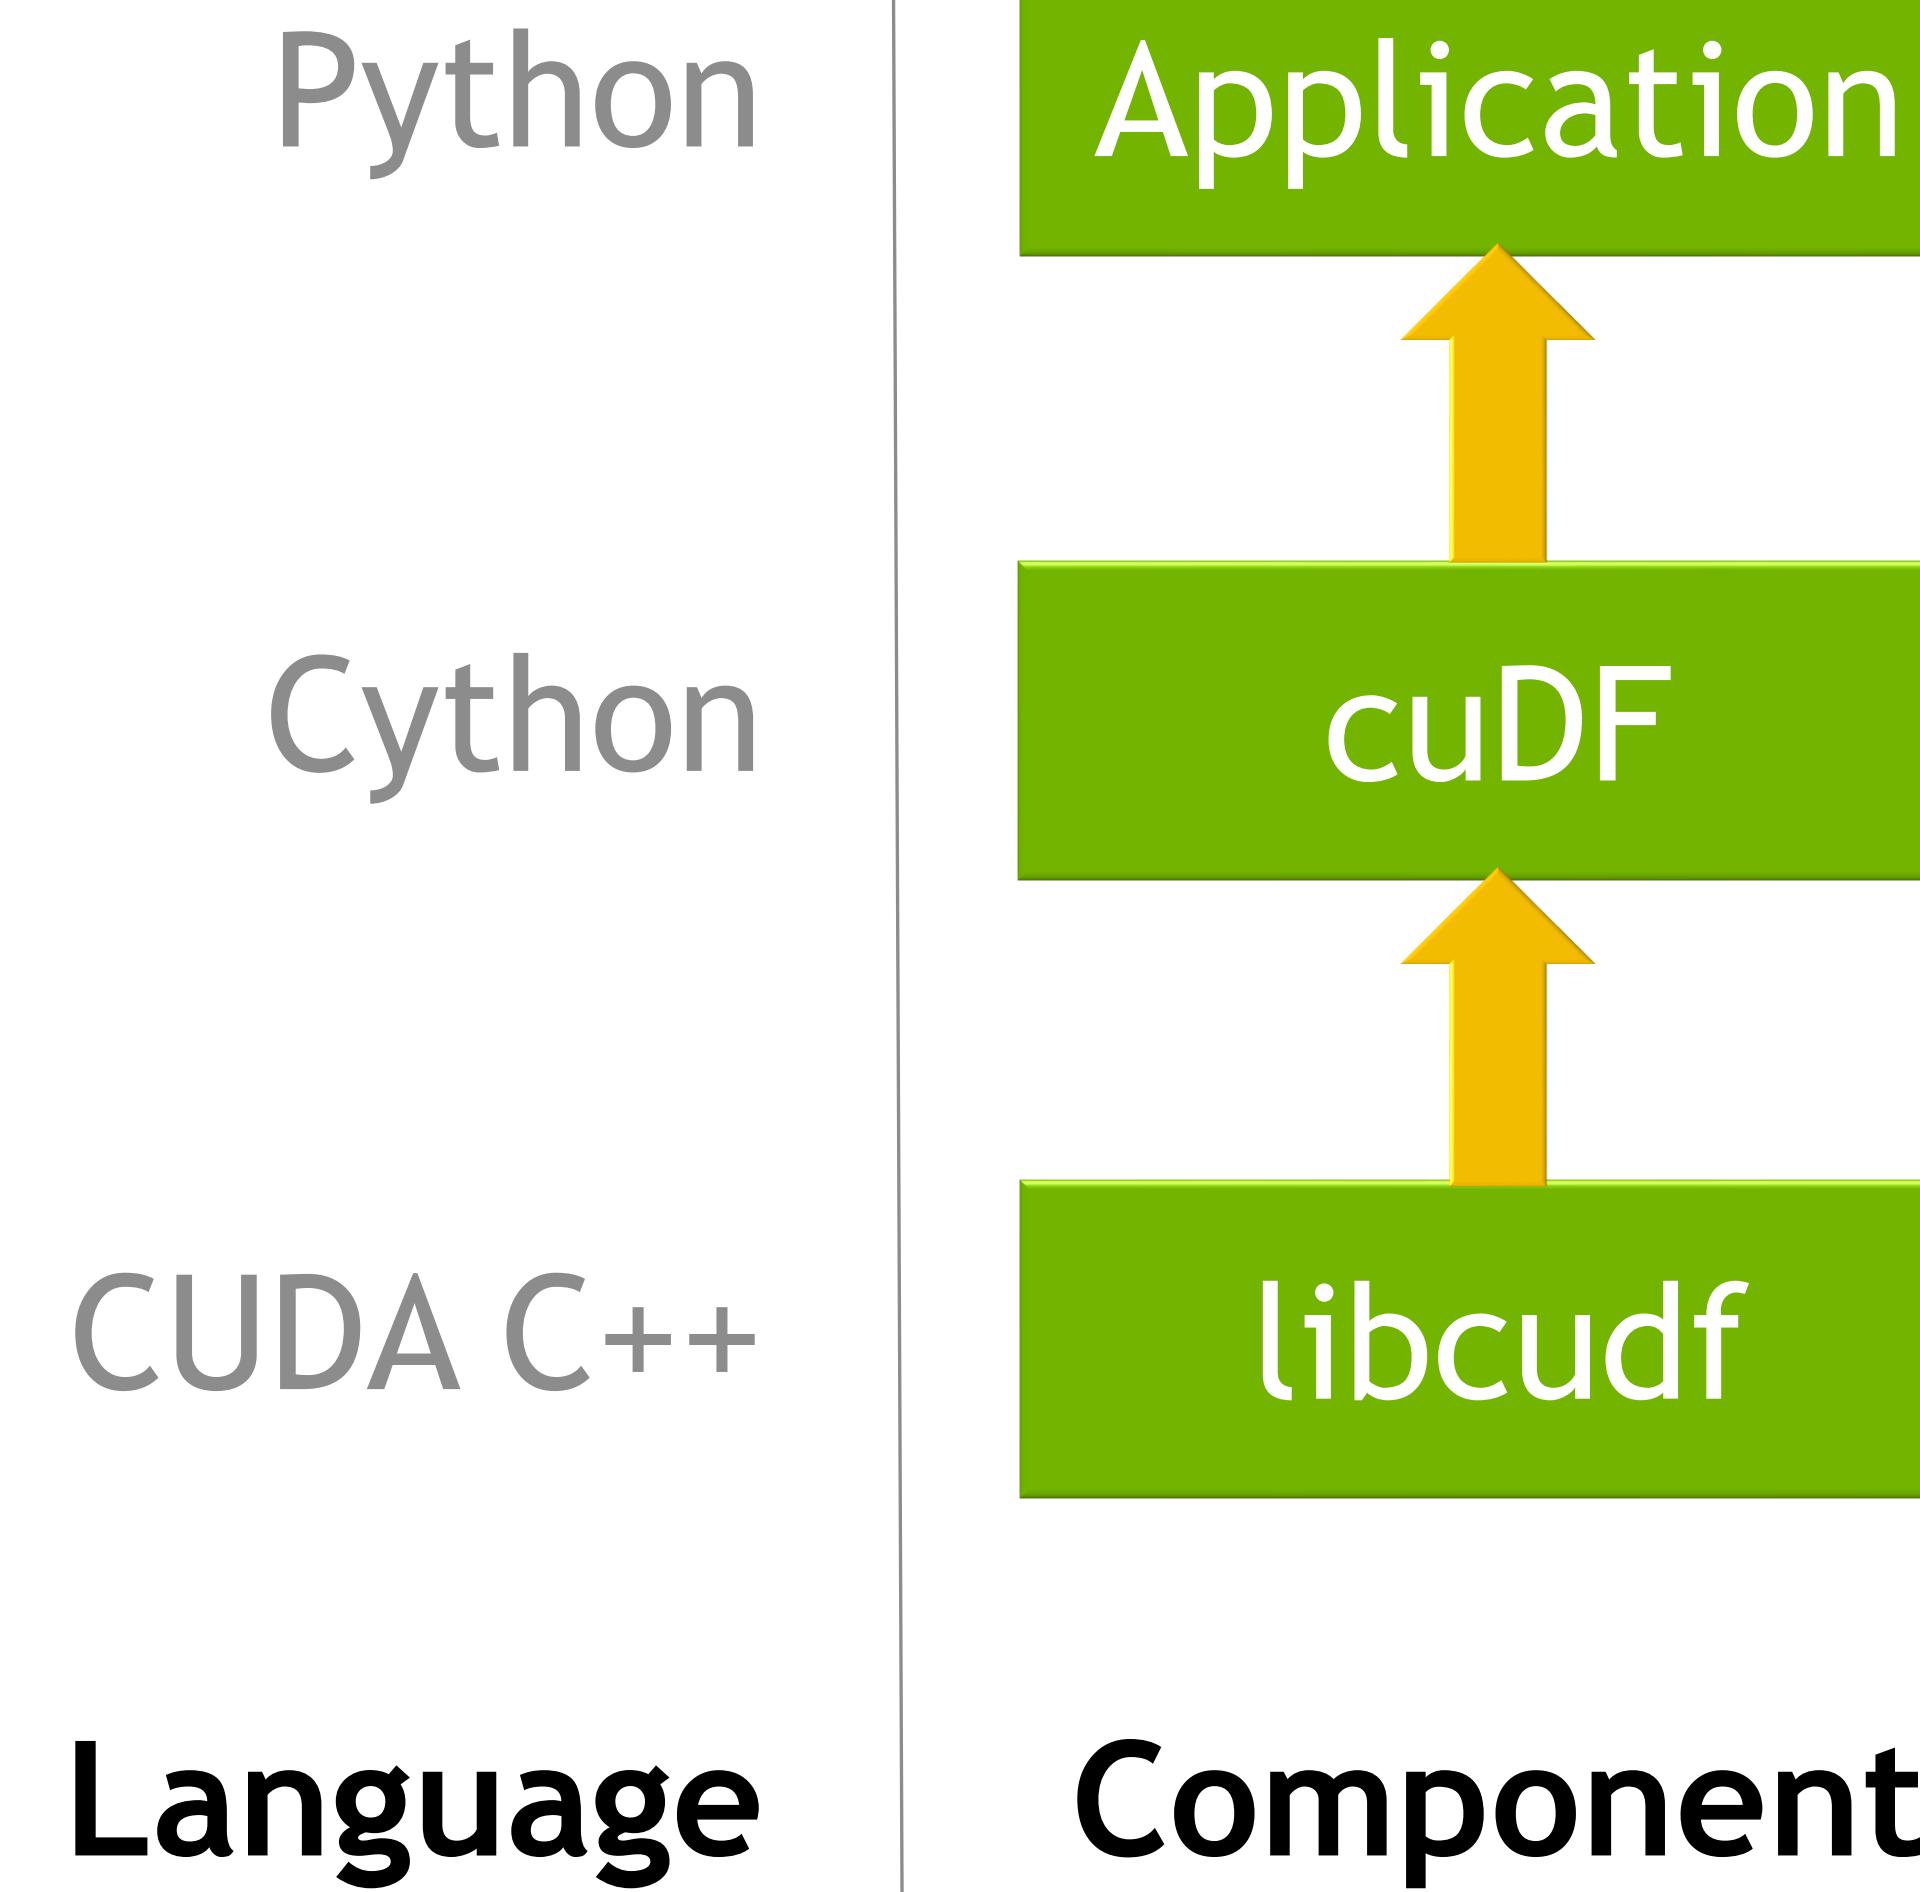
EXAMPLE 1 - CUDF

RAPIDS

- Pandas for GPUs, or:

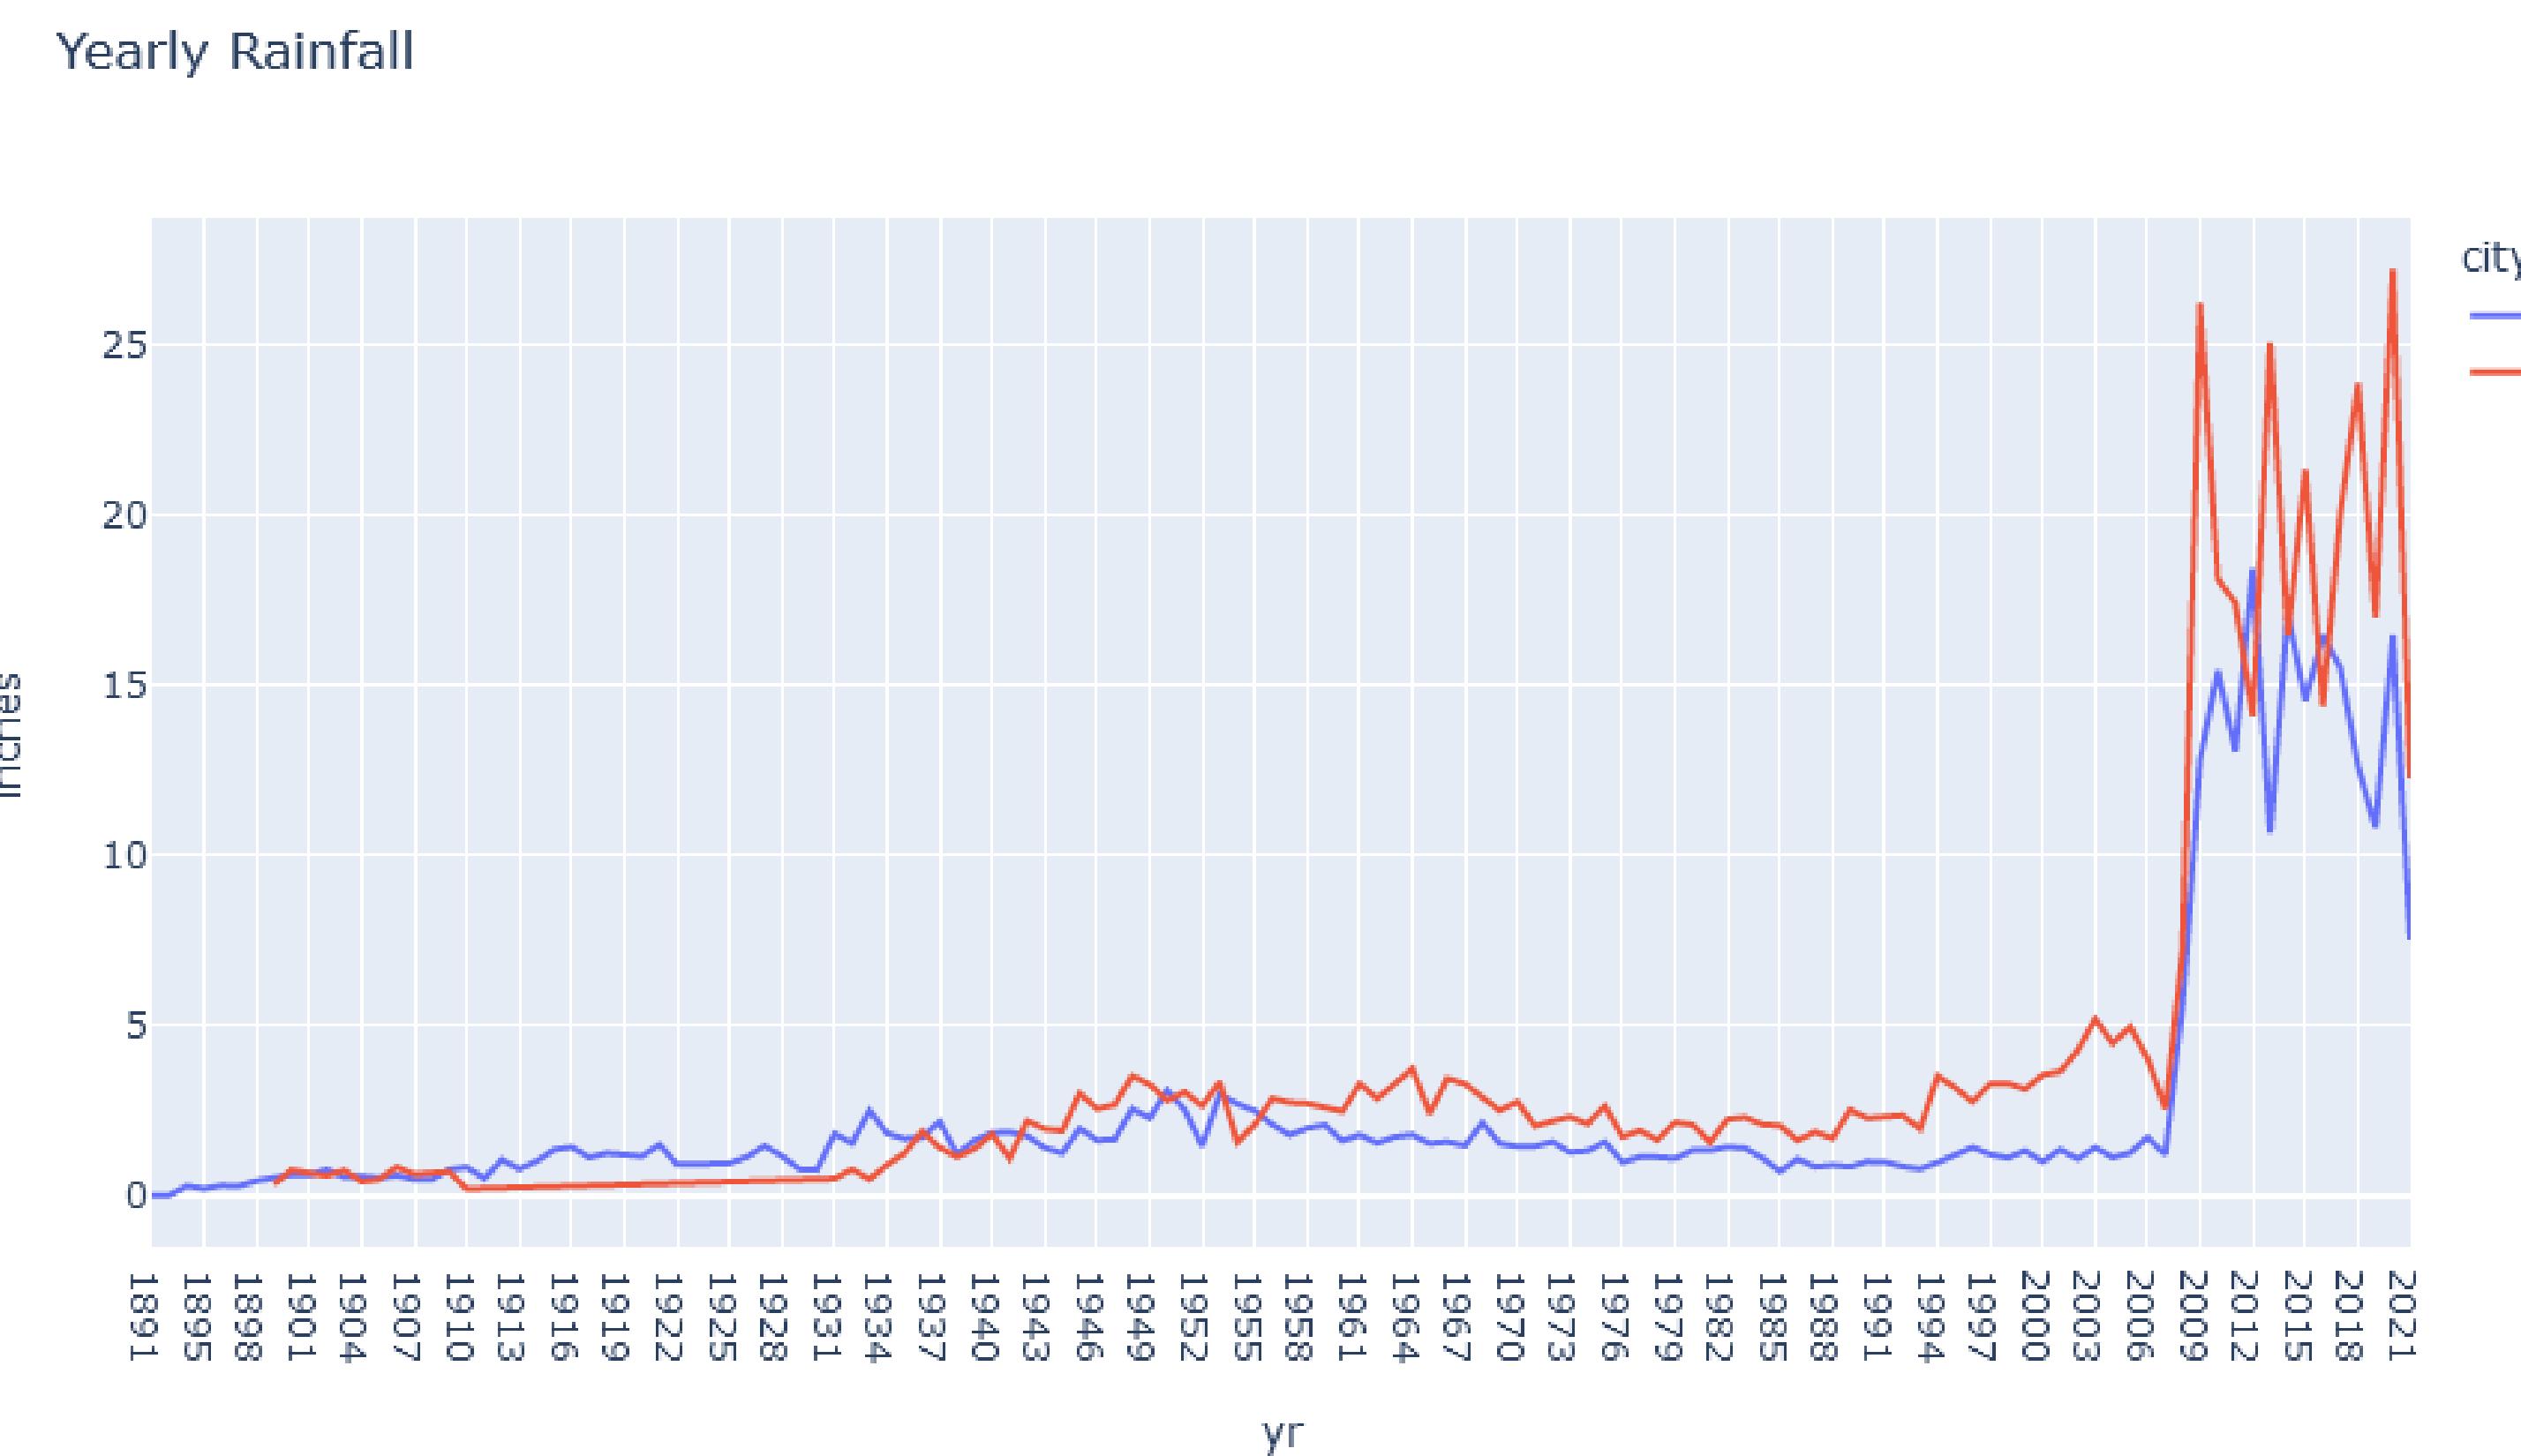
“cuDF is a Python GPU DataFrame library (built on the Apache Arrow columnar memory format) for loading, joining, aggregating, filtering, and otherwise manipulating tabular data using a DataFrame style API.”

```
# Defining a series:  
s = cudf.Series([1, 2, 3, None, 4])  
  
# Gives (2.5, 1.666666666666666)  
s.mean(), s.var()  
  
def add_ten(num):  
    return num + 10  
  
# Compiles add_ten() for CUDA GPU and runs it  
# Gives (11, 12, 13, <NA>, 14)  
s.applymap(add_ten)
```



EXAMPLE 1 - UDF PERFORMANCE IN CUDF AND DASK-SQL

- Example UDF from “[The Weather Notebook](#)”, [Dask-SQL for Data Exploration & Analysis](#)
- First presented in “[Accelerating Data Science: State of RAPIDS](#)” -John Zedlewski, Ben Zaitlen, Randy Gelhausen, GTC Fall 2021
- Query execution time **0.83s** on ~3M rows on 8-node dask cluster on DGX-1
- CPU execution time for comparison: **1.7s**

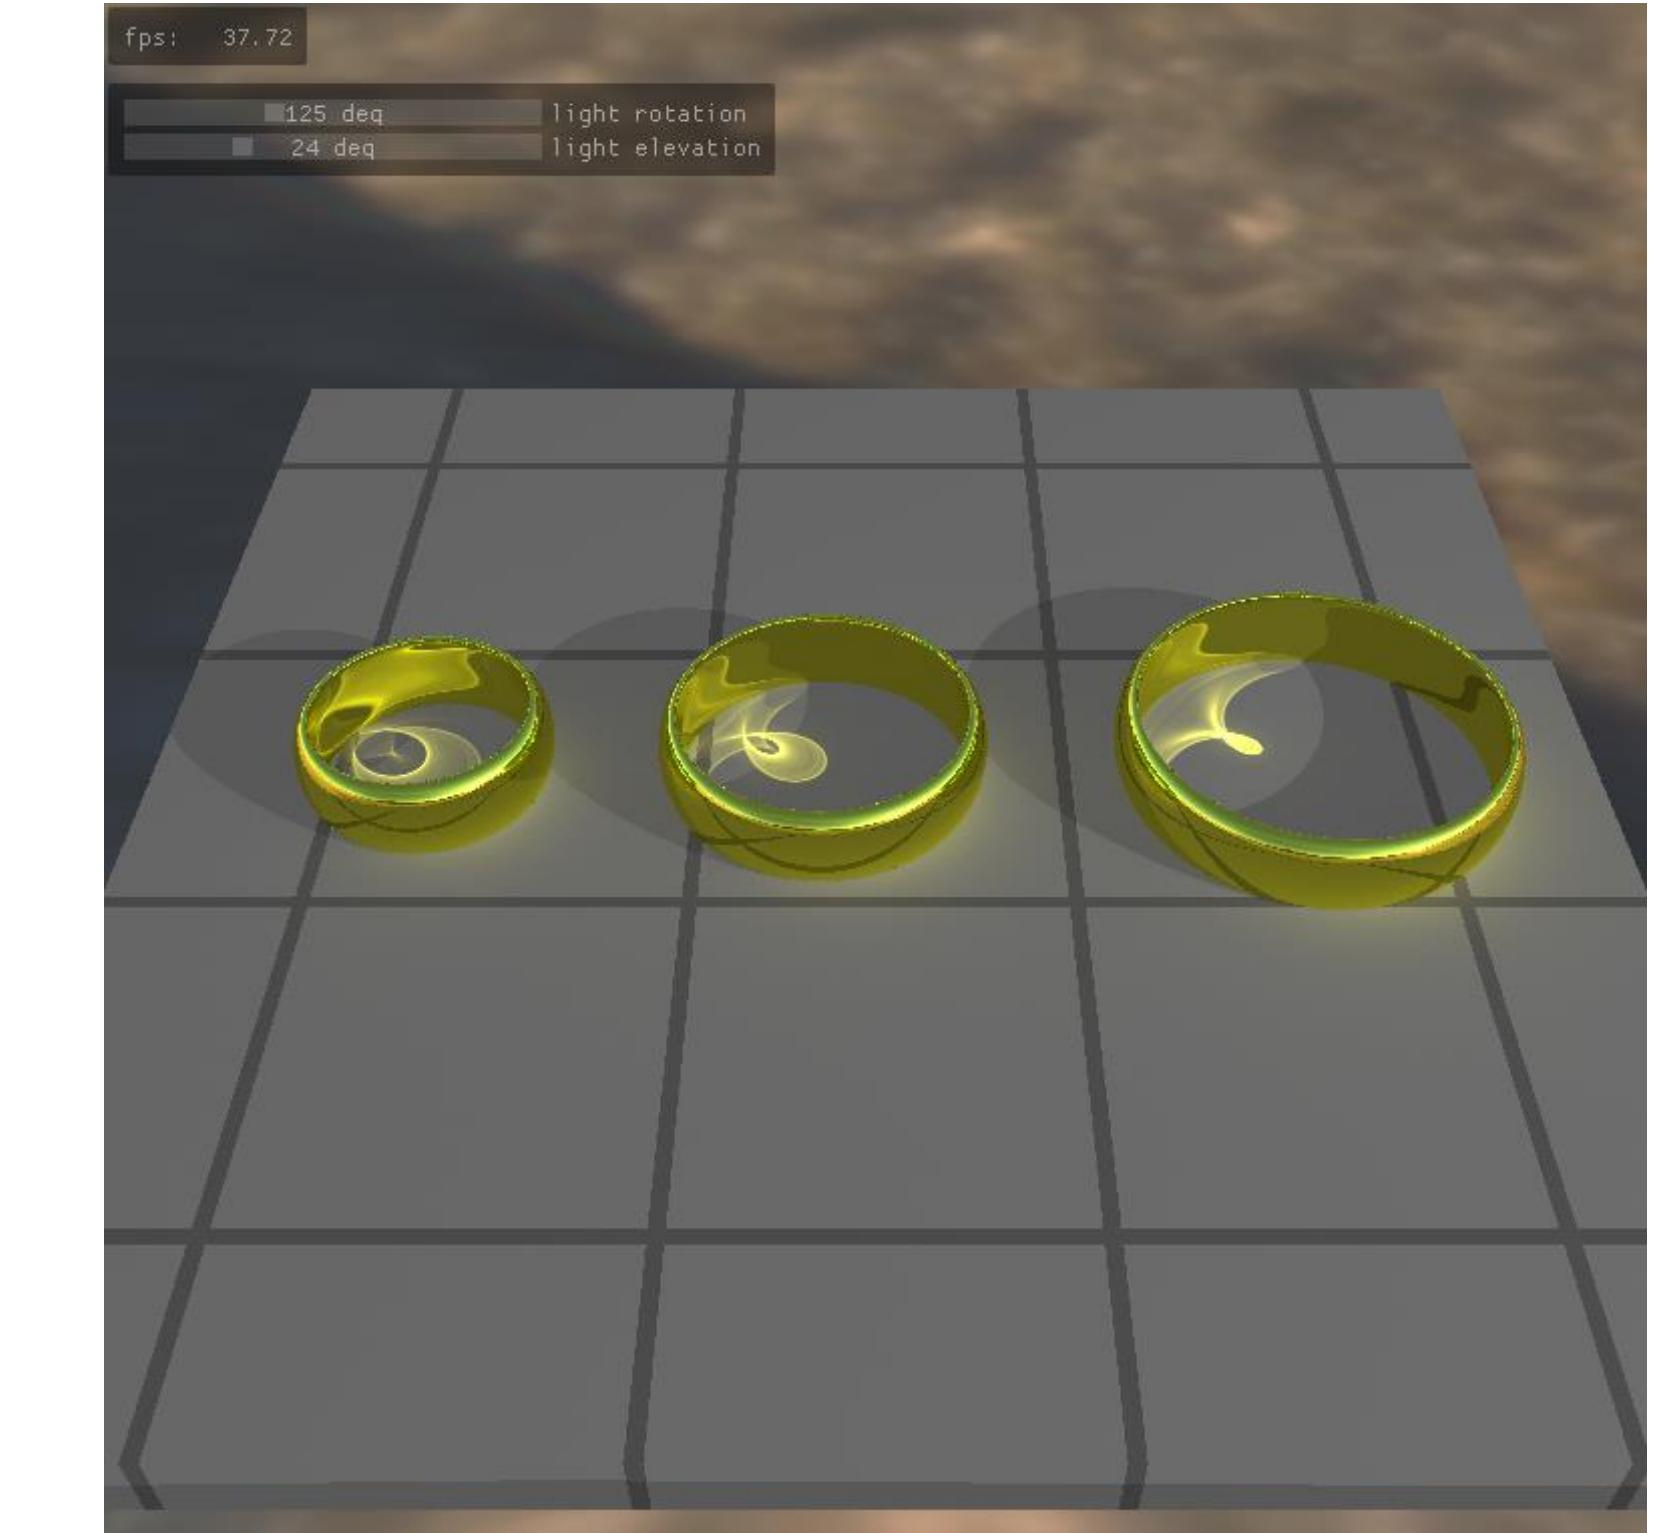
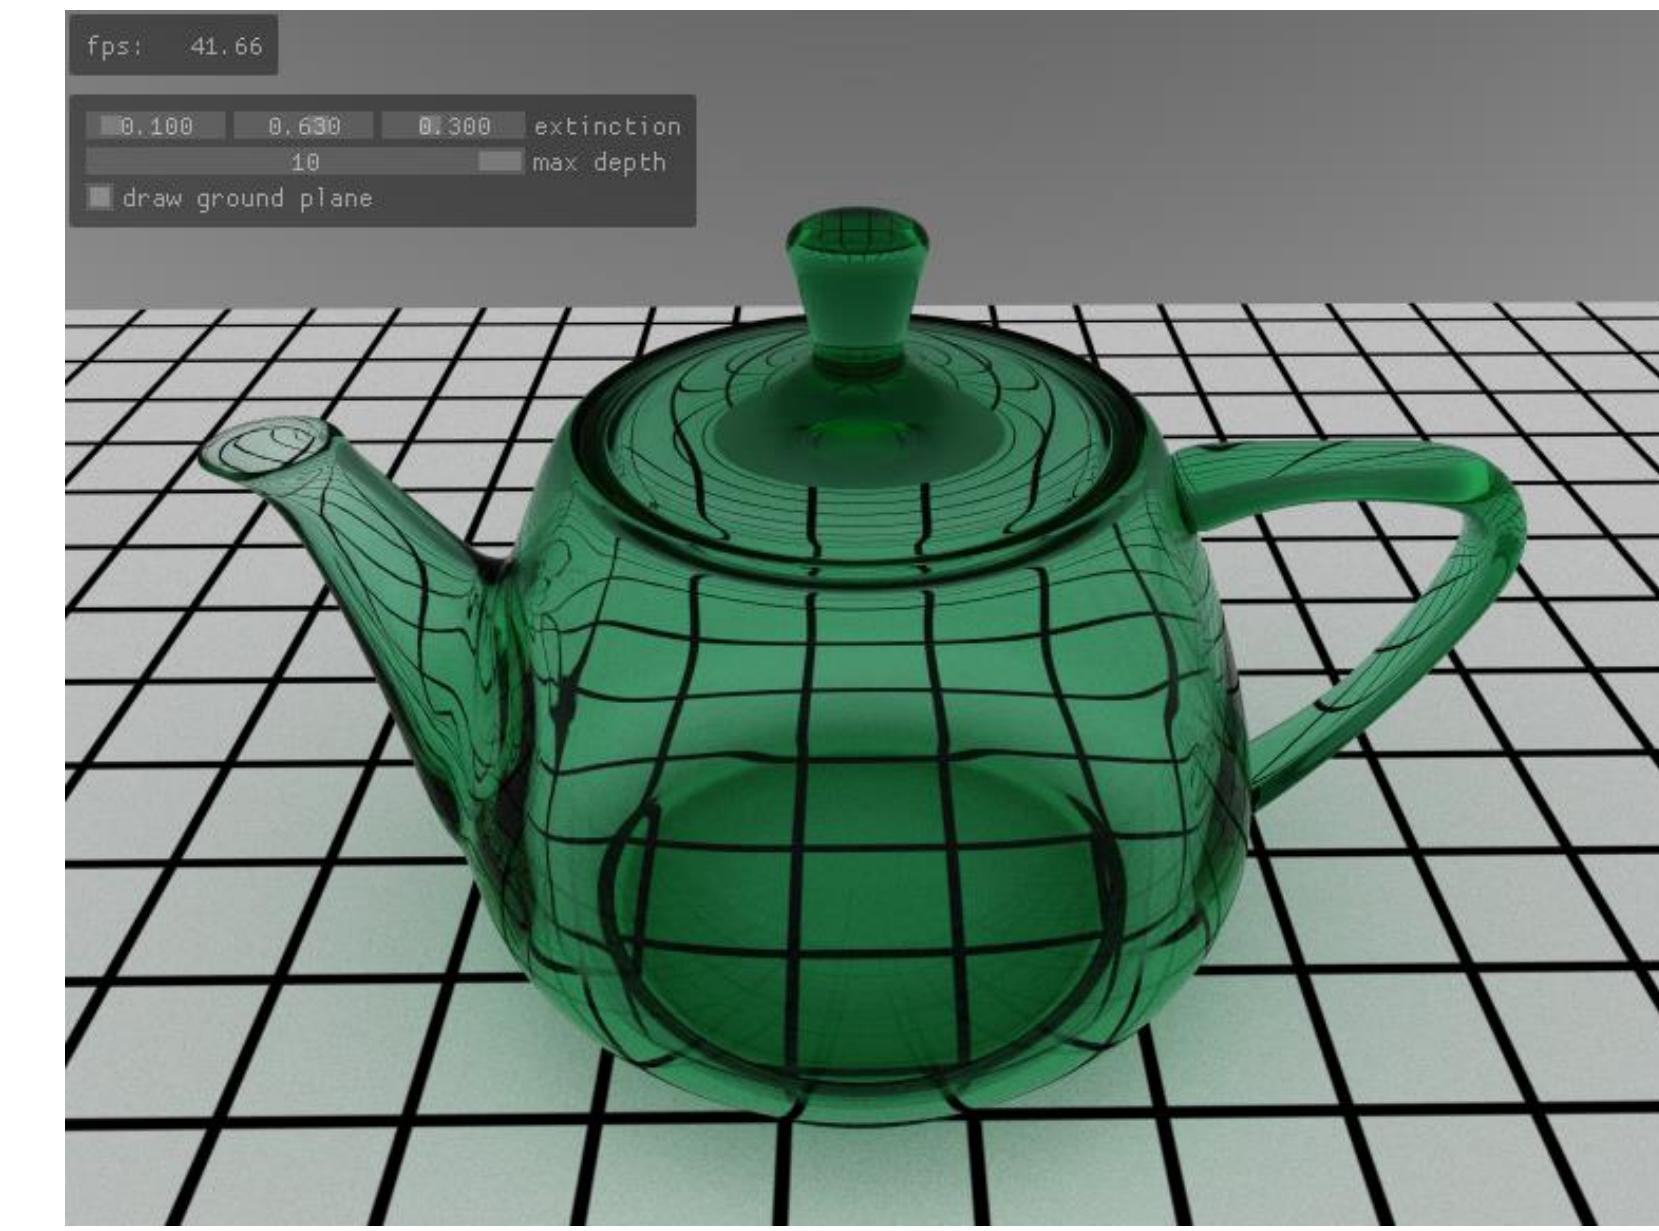


```
def haversine_dist(row, target_latitude, target_longitude):  
    x_1 = row["lat1"]  
    y_1 = row["lon1"]  
    x_2 = target_latitude  
    y_2 = target_longitude  
  
    x_1 = math.pi / 180 * x_1  
    y_1 = math.pi / 180 * y_1  
    x_2 = math.pi / 180 * x_2  
    y_2 = math.pi / 180 * y_2  
  
    dlon = y_2 - y_1  
    dlat = x_2 - x_1  
    a = (  
        math.sin(dlat / 2) ** 2  
        + math.cos(x_1) * math.cos(x_2) * math.sin(dlon / 2) ** 2  
    )  
  
    c = 2 * math.asin(math.sqrt(a))  
    r = 6371 # Radius of earth in kilometers  
  
    return c * r
```

EXAMPLE 2 - PYOPTIX

Images rendered by OptiX sample applications

- OptiX: optimal performance GPU-accelerated ray tracing
- PyOptiX: Python bindings for host-side functions, CUDA C/C++ kernels
- Numba + PyOptiX: write on-GPU raytracing kernels in Python



EXAMPLE 2 - PYOPTIX CUDA C/C++ KERNEL

```
static __forceinline__ __device__ void computeRay(uint3 idx, uint3 dim, float3& origin, float3& direction)
{
    const float3 U = params.cam_u;
    const float3 V = params.cam_v;
    const float3 W = params.cam_w;
    const float2 d = 2.0f * make_float2(
        static_cast<float>( idx.x ) / static_cast<float>( dim.x ),
        static_cast<float>( idx.y ) / static_cast<float>( dim.y )
    ) - 1.0f;

    origin     = params.cam_eye;
    direction = normalize( d.x * U + d.y * V + W );
}

extern "C" __global__ void __raygen_rg()
{
    // Lookup our location within the launch grid
    const uint3 idx = optixGetLaunchIndex();
    const uint3 dim = optixGetLaunchDimensions();

    // Map our launch idx to a screen location and create a ray from the camera
    // location through the screen
    float3 ray_origin, ray_direction;
    computeRay( make_uint3( idx.x, idx.y, 0 ), dim, ray_origin, ray_direction );
    // ...
}
```

EXAMPLE 2 - PYOPTIX PYTHON KERNEL WITH NUMBA

```
@cuda.jit(device=True, fast_math=True)
def computeRay(idx, dim):
    U = params.cam_u
    V = params.cam_v
    W = params.cam_w
    # Normalizing coordinates to [-1.0, 1.0]
    d = float32(2.0) * make_float2(
        float32(idx.x) / float32(dim.x), float32(idx.y) / float32(dim.y)
    ) - float32(1.0)

    origin = params.cam_eye
    direction = normalize(d.x * U + d.y * V + W)
    return origin, direction

def __raygen_rg():
    # Lookup our location within the launch grid
    idx = optix.GetLaunchIndex()
    dim = optix.GetLaunchDimensions()

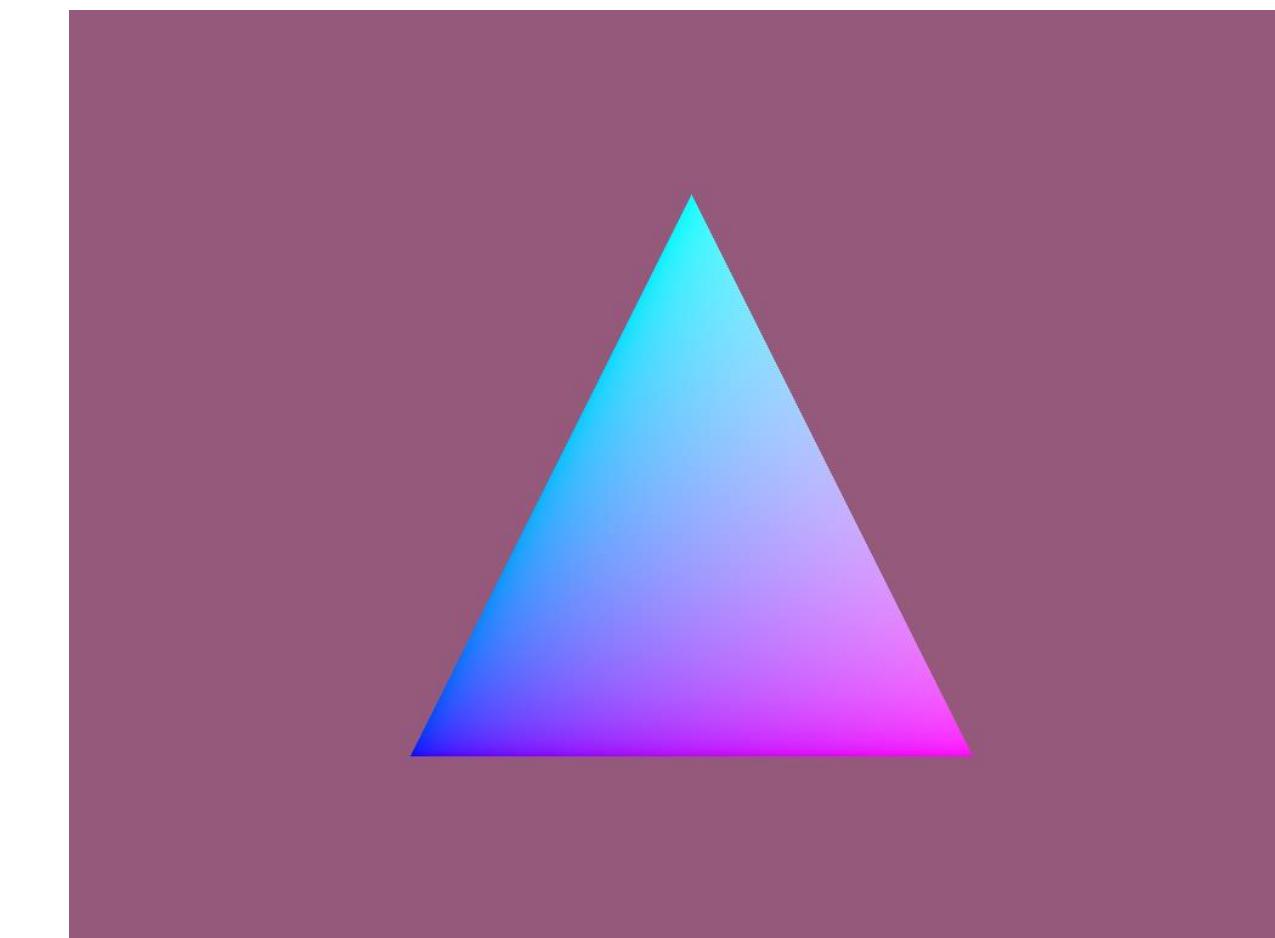
    # Map our launch idx to a screen location and create a ray from the camera
    # location through the screen
    ray_origin, ray_direction = computeRay(make_uint3(idx.x, idx.y, 0), dim)

    # ...
```

EXAMPLE 2 - PYOPTIX RAYGEN KERNEL PERFORMANCE

- Kernel execution time measured with Nsight Compute:

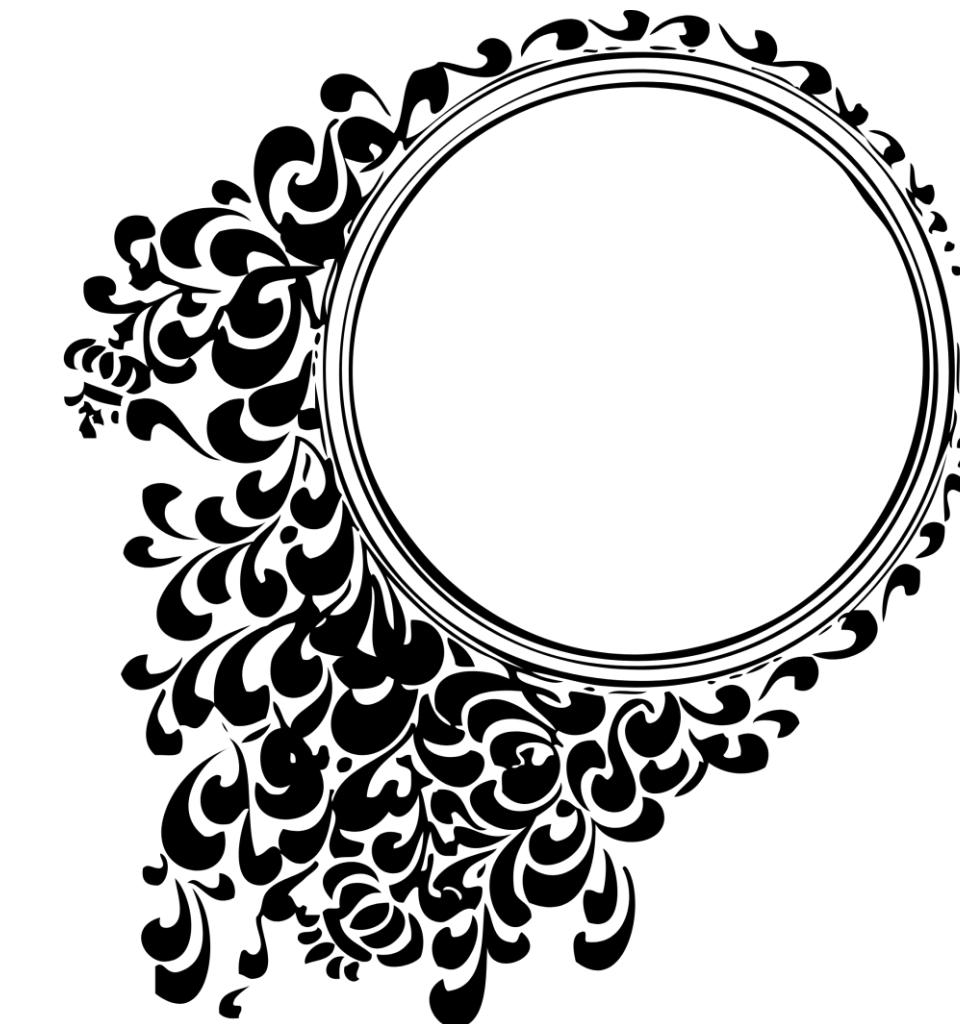
Language	Kernel execution time (cycles)	% of baseline
C++	94,172	100.0
Python	106,776	113.3



Triangle rendered by example kernel

- Further optimization: force inlining, fastmath flags
 - Target: Numba 0.56 (June / July)

DEMO APPLICATION - FILIGREE



- **Filigree:** an image processing library
 - Written in CUDA C++
 - Uses ImageMagick for file I/O and data structures
 - Python Bindings + API

```
__global__ void to_greyscale_kernel(
    Magick::PixelPacket *pixels,
    size_t width, size_t height) {
    size_t x = blockIdx.x * blockDim.x + threadIdx.x;
    size_t y = blockIdx.y * blockDim.y + threadIdx.y;

    if ((x < width) && (y < height)) {
        Magick::PixelPacket *pixel = &pixels[y * width + x];
        float weighted = (0.299f * pixel->red +
                           0.587f * pixel->green +
                           0.114 * pixel->blue);
        pixel->red = (Magick::Quantum)weighted;
        pixel->green = (Magick::Quantum)weighted;
        pixel->blue = (Magick::Quantum)weighted;
    }
}
```

- Full source, completed example:
 - <https://github.com/gmarkall/numba-accelerated-udfs>

PYTHON “HOST” API

```
from filigree import Image  
image = Image(path)  
image.to_pillow()
```

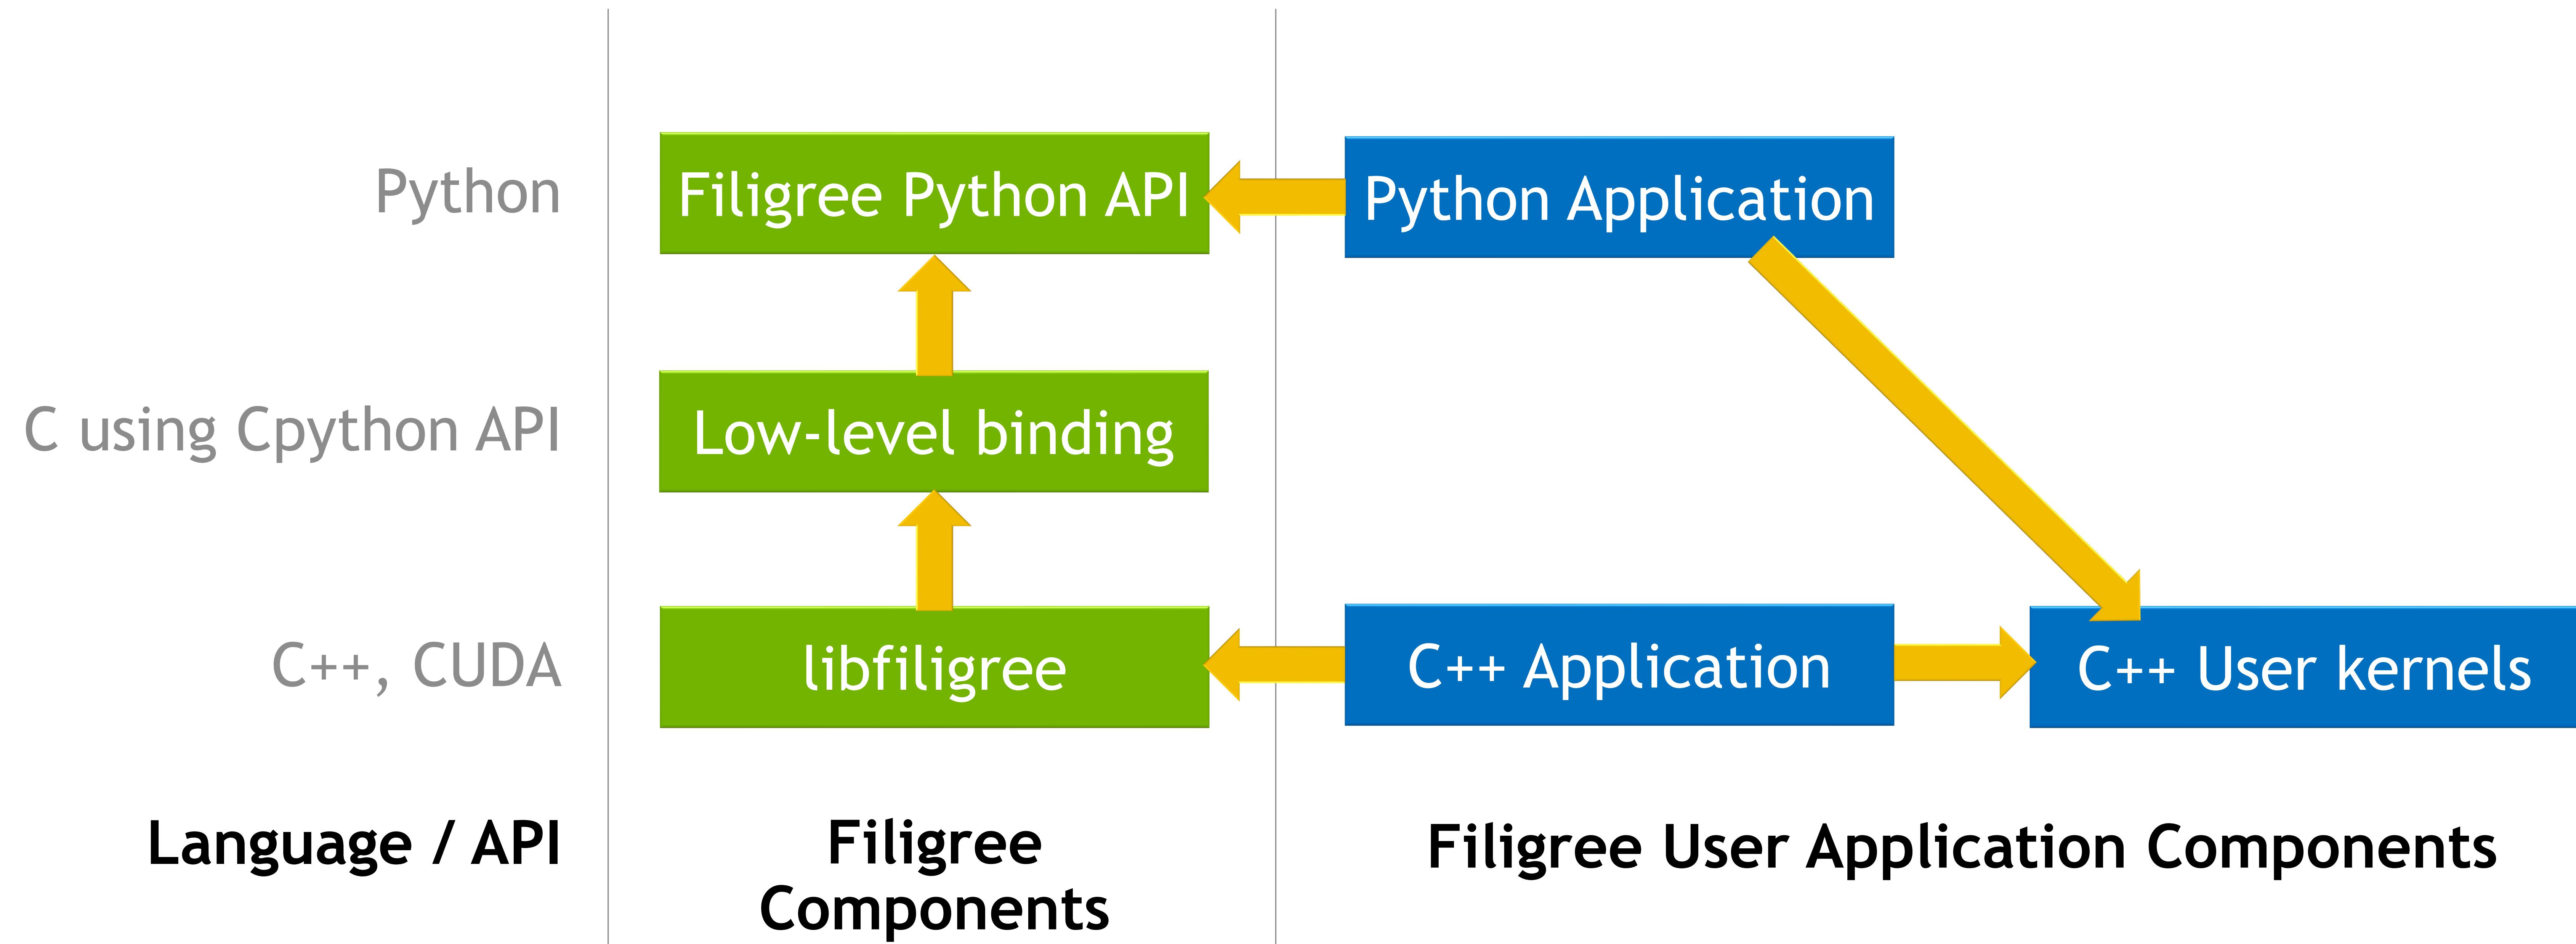


```
image.apply_weighted_greyscale()
```

```
image.to_pillow()
```



FILIGREE COMPONENTS



AIM: SUPPORT THIS USE CASE

- User-defined function, written in Python, applied from Python
- Compiled to run on GPU and operate directly on Filigree data
- No user-visible C / C++ / CUDA / low-level API

```
image = Image(path)

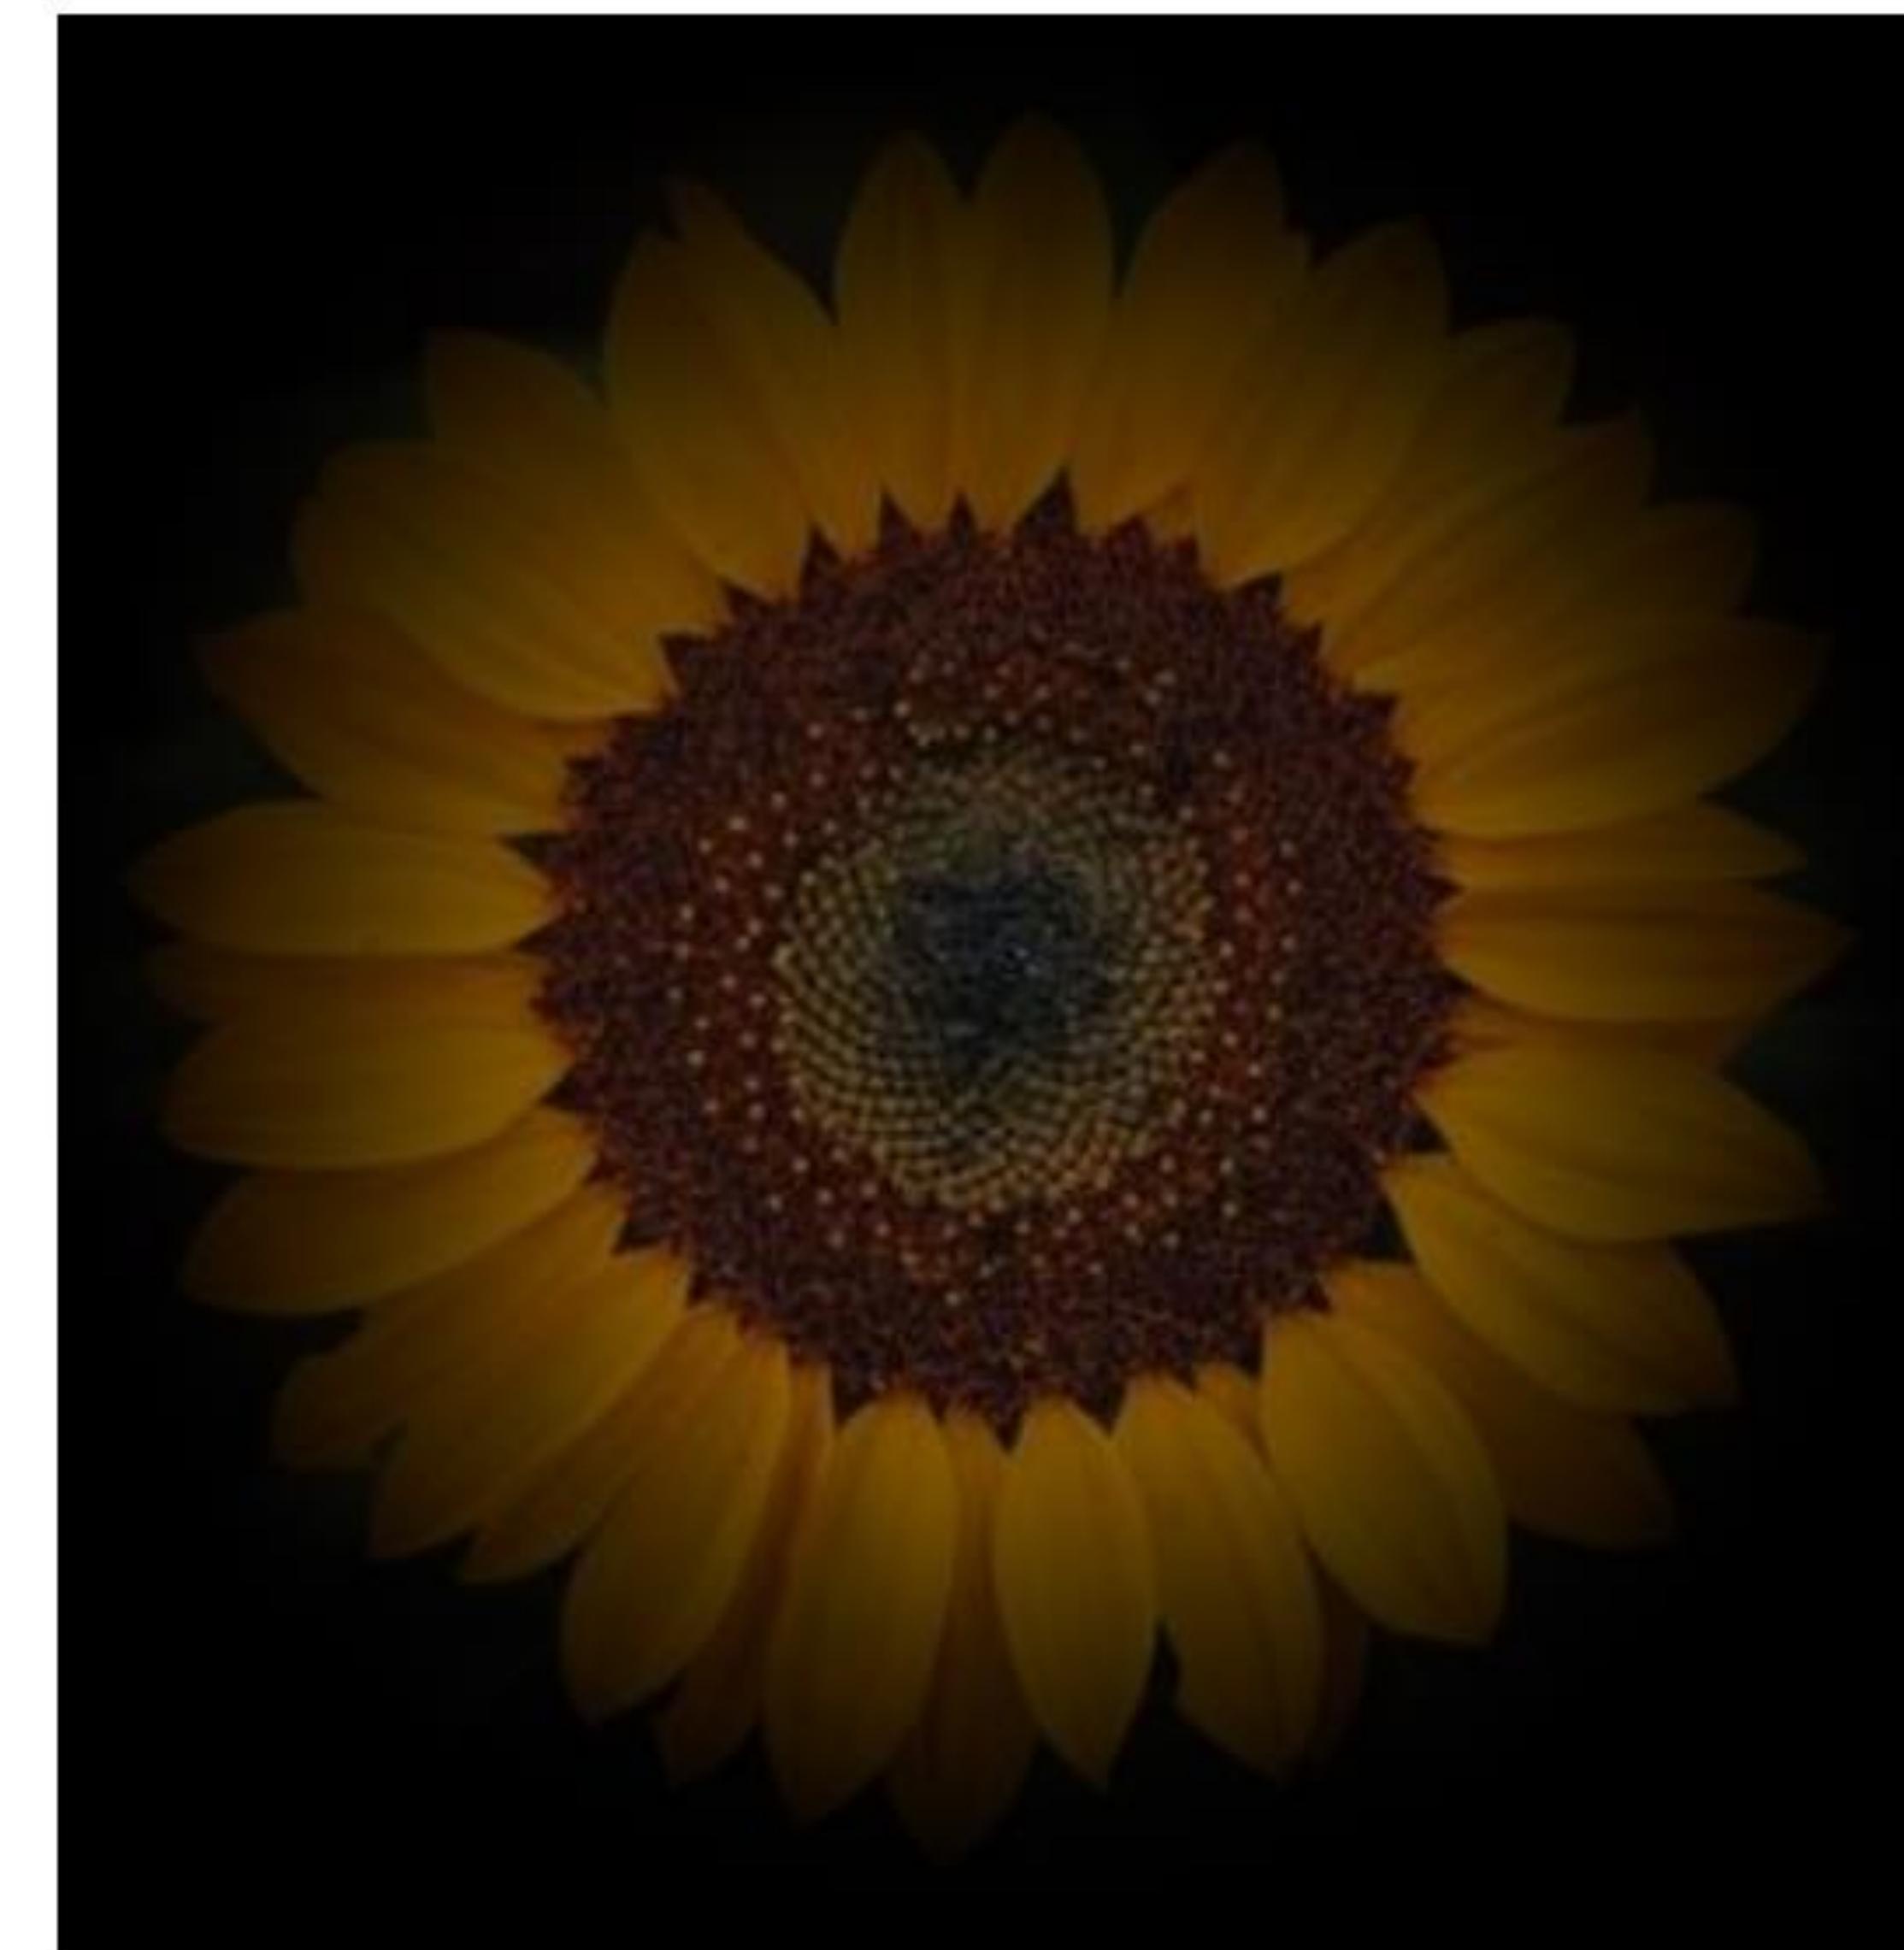
center_x = image.width / 2
center_y = image.height // 2
radius = max(center_x, center_y)

def highlight_center(pixel, x, y):
    distance_from_center = math.sqrt(
        abs(x - center_x)**2 + abs(y - center_y)**2)

    # Weight pixels according to distance from centre
    w = 1 - (distance_from_center / radius)

    # Apply weight to each RGB component
    return pixel.r * w, pixel.g * w, pixel.b * w, pixel.a

image.apply_located_pixel_udf(highlight_center)
image.to_pillow()
```



TECHNICAL FOUNDATION

Numba - a Python JIT Compiler



A BRIEF INTRODUCTION TO NUMBA

■ What is Numba?

- A Just-in-time (JIT) compiler for Python functions.
- Opt-in: Numba only compiles the functions you specify
- Focused on array-oriented and numerical code
- Trade-off: subset of Python for better performance
- Alternative to native code, e.g. C / Fortran / Cython / CUDA C/C++
- Targets:
 - CPUs: x86, PPC, ARMv7 / v8
 - GPU: CUDA

■ CUDA Target:

- Compiler
- Driver bindings
- Device array library

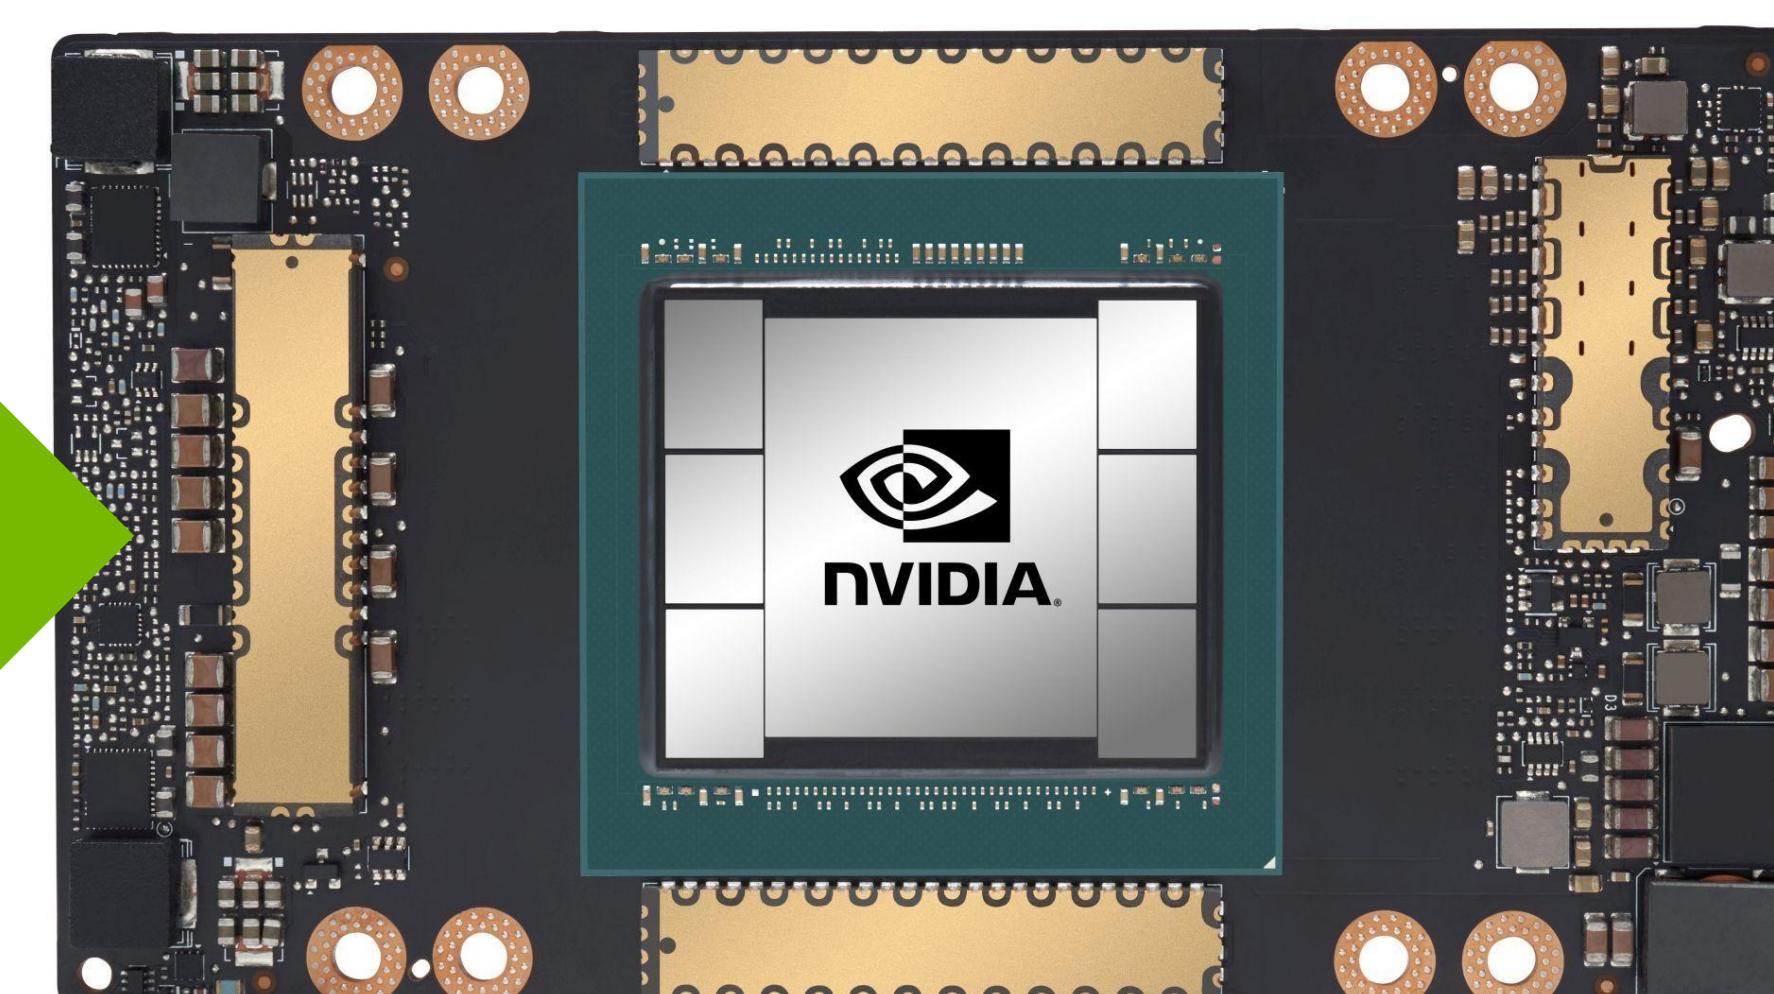
```
@cuda.jit
def increment_a_2D_array(an_array):
    x, y = cuda.grid(2)
    if x < an_array.shape[0] and y < an_array.shape[1]:
        an_array[x, y] += 1
```



CUDA Python Ecosystem

RAPIDS Numba
CuPy PyTorch
DALI

+ more... See: [CUDA Array Interface documentation](#)



“STANDALONE” NUMBA USAGE

```
from numba import cuda

# Define a kernel that is compiled for CUDA
@cuda.jit
def vector_add(r, x, y):
    start = cuda.grid(1)
    step = cuda.gridsize(1)
    stop = len(r)

    for i in range(start, stop, step):
        r[i] = x[i] + y[i]

# Allocate some arrays on the device and copy data
N = 2 ** 10
x = cuda.to_device(np.arange(N))
y = cuda.to_device(np.arange(N) * 2)
r = cuda.device_array_like(x)

# Configure and launch kernel
block_dim = 256
grid_dim = (len(x) // block_dim) + 1
vector_add[grid_dim, block_dim](r, x, y)

# Copy result back from the device
result = r.copy_to_host()
```

■ Numba + CUDA Tutorial:

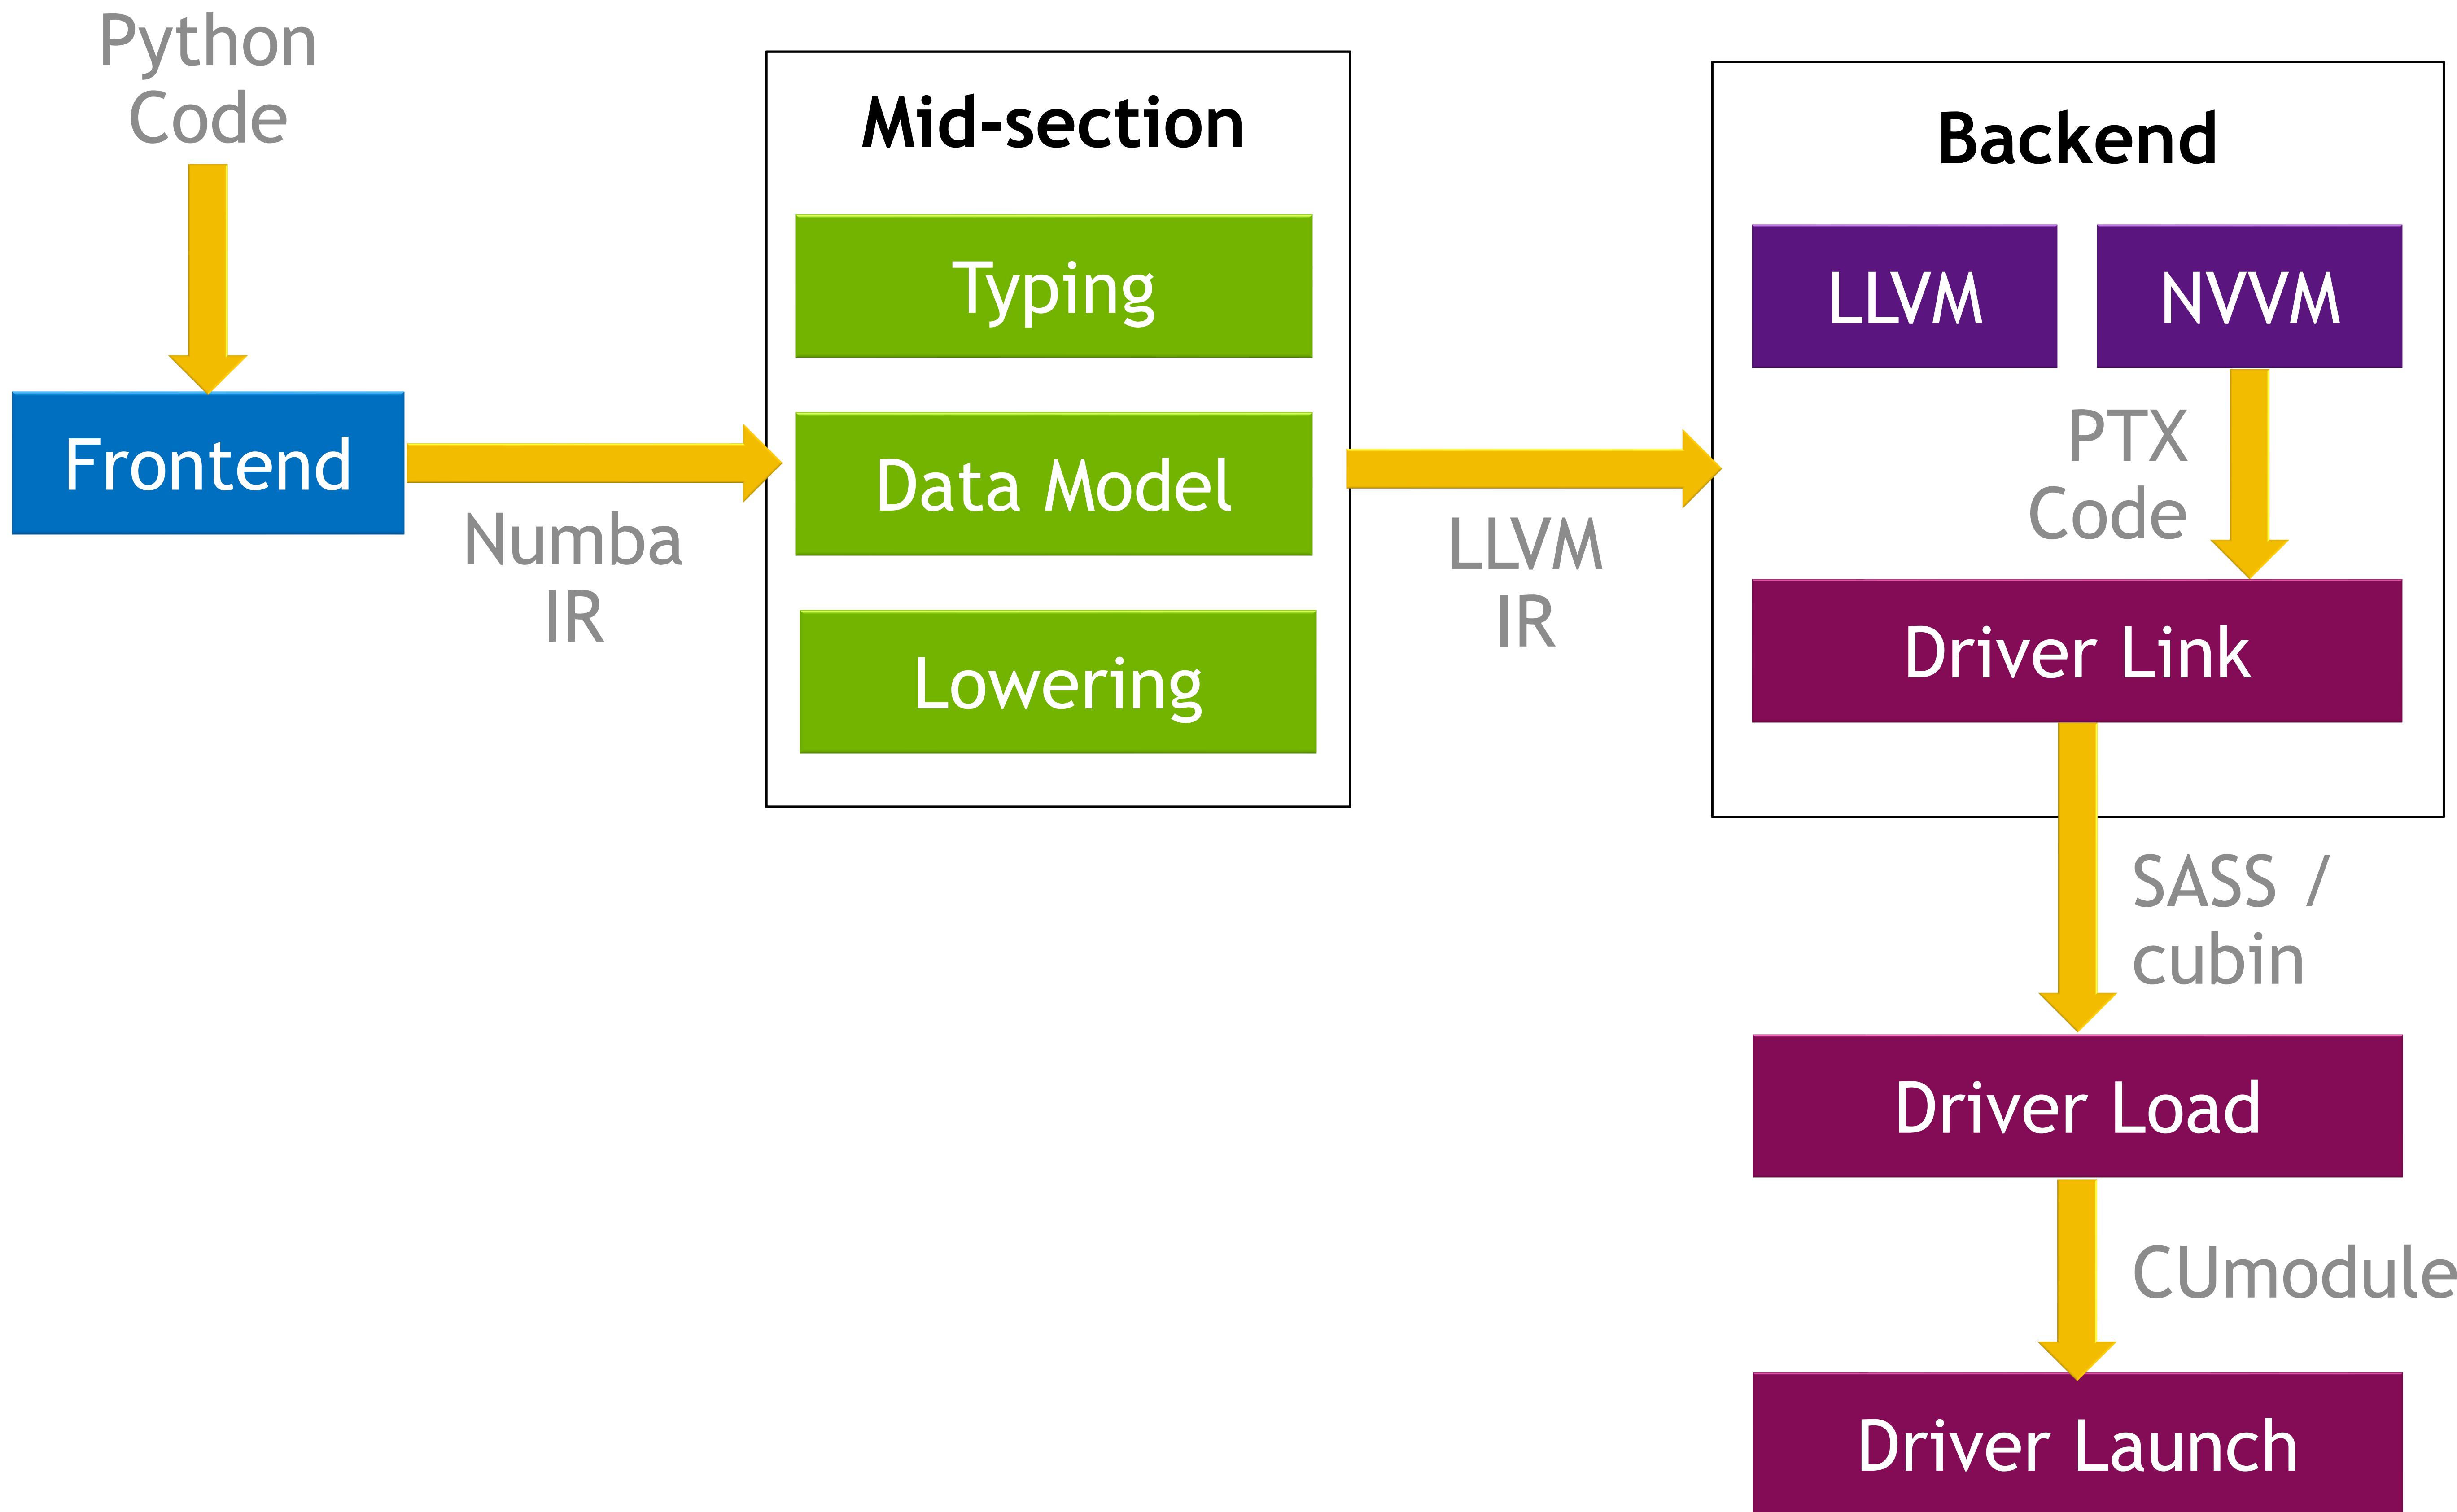
<https://github.com/numba/nvidia-cuda-tutorial>

- **Session 1:** An introduction to Numba and CUDA Python
- **Session 2:** Typing
- **Session 3:** Porting strategies, performance, interoperability, debugging
- **Session 4:** Extending Numba
- **Session 5:** Memory Management

WHO IS USING NUMBA?

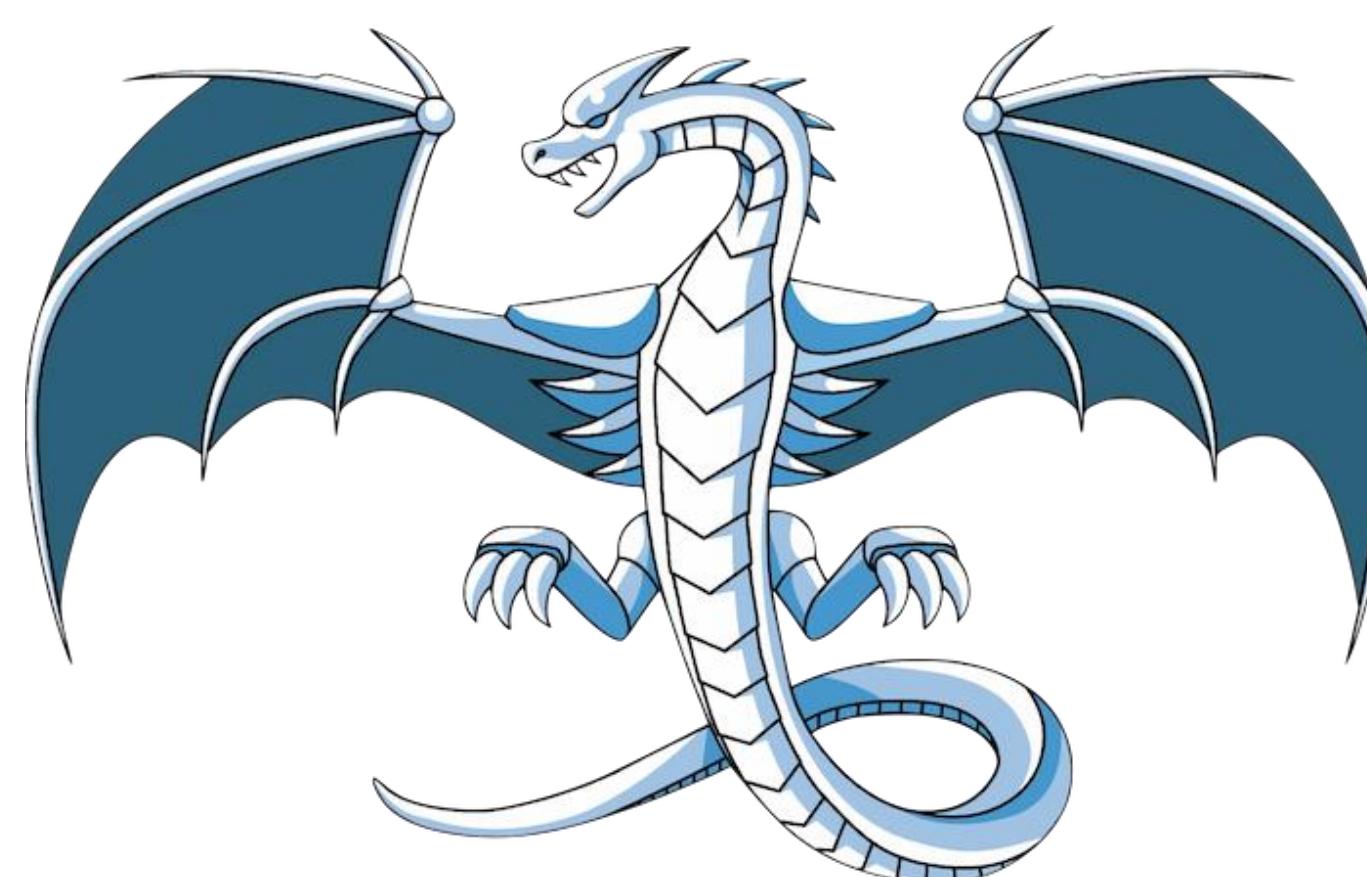
- PyPI: 250,000 / Conda 16,000+ downloads per day
- Github:
 -  48,000 dependent repositories
 -  7,300 stars
 -  879 forks
 -  205 watchers
- Random sample of applications using the CUDA target:
 - Poliastro (astrodynamics)
 - FBPIC (CUDA-accelerated plasma physics)
 - UMAP (manifold learning)
 - RAPIDS (data science)
 - [Talks on more applications in the Numba documentation](#)
 - and <https://github.com/gmarkall/numba-cuda-users>

NUMBA PIPELINE



LLVM

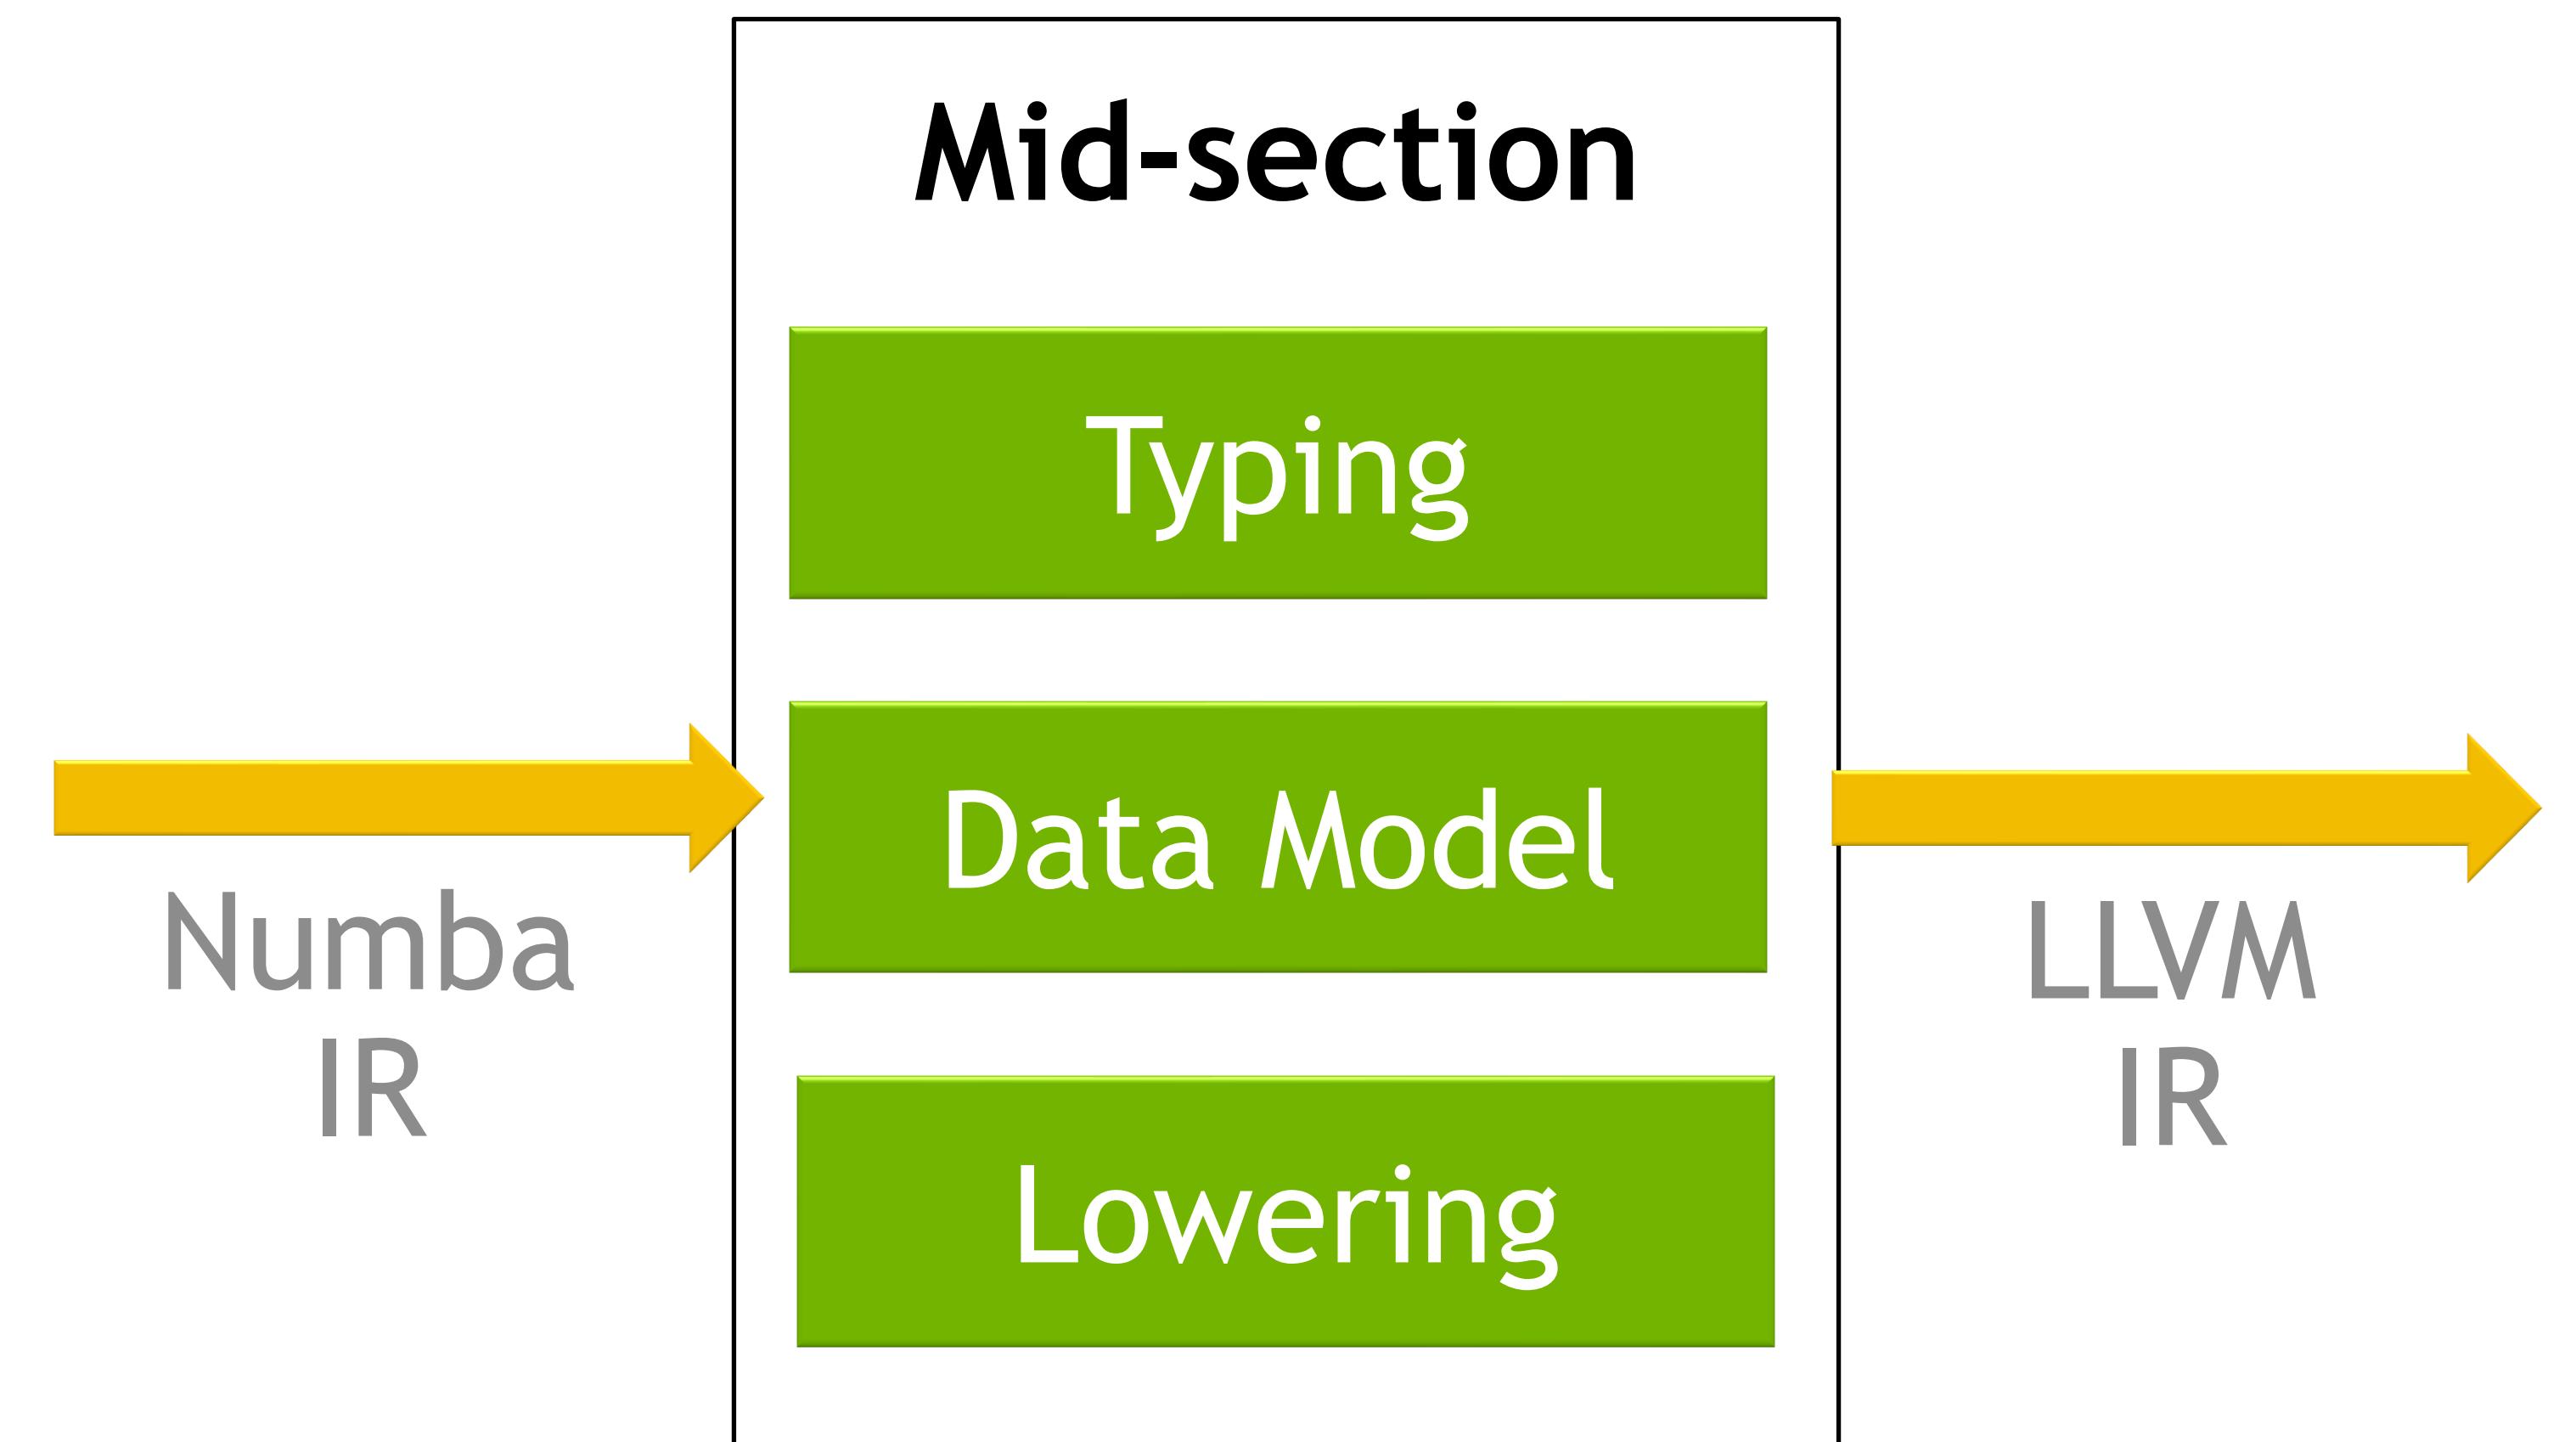
- “The LLVM Project is a collection of modular and reusable compiler and toolchain technologies.” - <https://llvm.org>



The LLVM Wyvern logo

UDF-RELATED COMPONENTS

- **Typing (adds type info):**
 - Teach Numba to recognise our library's types,
 - and functions operating on those types
- **Data model (Maps Numba -> LLVM types):**
 - Add mappings for our types / data structures
- **Lowering (Numba IR -> LLVM IR):**
 - Add implementations of operations on our library data structures



TYPING

- **Type Inference:** types determined by propagating type information
 - **Function Typing:** “*For these argument types, what type is returned?*”
 - **Data flow:**
 - “*Where did the inputs come from?*”
 - “*Where is the output / return going?*”
- **“No regrets”:** typing does not backtrack
- Simple example:

```
def f(a, b):          # a:= float32, b:= int32
    c = np.log(b)      # c:= float64
    d = a + c          # d = float64
    return d           # return := float64
```

TYPE UNIFICATION (1)

```
def select(a, b, c): # a := float32, b := float32, c := bool
    if c:
        ret = a      # ret := float32
    else:
        ret = b      # ret := float32
    return ret       # return := {float32, float32}
                    #           => float32
```

TYPE UNIFICATION (2)

```
def select(a, b, c): # a := float32, b := float64, c := bool
    if c:
        ret = a      # ret := float32
    else:
        ret = b      # ret := float64
    return ret       # return := {float32, float64}
                    #           => float64
```

TYPE UNIFICATION (3)

```
def select(a, b, c): # a := tuple(int32, int32), b := float64, c := bool
    if c:
        ret = a      # ret := tuple(int32, int32)
    else:
        ret = b      # ret := float64
    return ret       # return := {tuple(int32, int32), float64}
                    #           => XXX!!!
```

numba.core.errors.TypingError: Failed in nopython mode pipeline (step: nopython frontend)
Can't unify return type from the following types: UniTuple(int32 x 2), float64

Unification rules:

- Internal to Numba
- Can be extended

HOW TO ADD TYPING - CONCRETE TYPING

- **Concrete:** a table of *specific* cases

- Examples from within Numba:

```
class Math_pow(ConcreteTemplate):  
    key = math.pow  
    cases = [  
        signature(types.float64, types.float64, types.int64),  
        signature(types.float64, types.float64, types.uint64),  
        signature(types.float32, types.float32, types.float32),  
        signature(types.float64, types.float64, types.float64),  
    ]
```

```
class Cuda_popc(ConcreteTemplate):  
    key = cuda.popc  
    cases = [  
        signature(types.int8, types.int8),  
        signature(types.int16, types.int16),  
        signature(types.int32, types.int32),  
        signature(types.int64, types.int64),  
        signature(types.uint8, types.uint8),  
        signature(types.uint16, types.uint16),  
        signature(types.uint32, types.uint32),  
        signature(types.uint64, types.uint64),  
    ]
```

HOW TO ADD TYPING - GENERIC / ABSTRACT TYPING

- Write a function that accepts argument types and computes return type:

```
# Simplified typing pseudo code:
def cuda_grid_typer(ndim):
    require_integer_literal(ndim)
    value = ndim.literal_value

    if value == 1:
        return int32
    elif value == 2:
        return tuple(int32, int32)
    elif value == 3:
        return tuple(int32, int32, int32)
    else:
        raise Error("Grid can only be 1D, 2D, or 3D")
```

DATA MODELS

- Connect the frontend to the backend
 - Mapping Numba types -> LLVM Types
- LLVM Type System:
 - Void (`void`)
 - Scalars
 - Integers: `i32`, `i64`, ...
 - Floating point: `float`, `double`, ...
 - Aggregates
 - Arrays: `[40 x i32]`
 - Structures: `{i32, i32, i64}`
 - Pointers (`void*`, `i32*`, `[40 x i32]*`, `float***`, ...)
- LLVM IR Reference manual:
<https://llvm.org/docs/LangRef.html#type-system>

DATA MODEL EXAMPLES

Python Value	Numba Type	LLVM Type used in Numba
3 (int)	int64	i64
3.14 (float)	float64	double
(1, 2, 3)	UniTuple(int64, 3)	[3 x i64]
(1, 2.5)	Tuple(int64, float64)	{i64, double}
np.array([1, 2], dtype=np.int32)	array(int64, 1d, C)	{i8*, i8*, i64, i64, i32*, [1 x i64], [1 x i64]}
“Hello”	unicode_type	{i8*, i64, i32, i32, i64, i8*, i8*}

TUPLE REPRESENTATIONS

Layer	Representation
Python	(1, 2.5)
Numba	<code>Tuple(int64, float64)</code>
C / C++	<code>struct {</code> <code> uint64_t v1;</code> <code> double v2;</code> <code>};</code>
LLVM	<code>{</code> <code> i64, ; v1</code> <code> double ; v2</code> <code>}</code>

ARRAY REPRESENTATIONS

Layer	Representation
Python	<code>np.array([1, 2], dtype=np.int32)</code>
Numba	<code>array(int64, 1d, C)</code>
C / C++	<pre>struct { void *meminfo; PyObject *parent; npy_intp nitems; npy_intp itemsize; void *data; npy_intp shape_and_strides[]; };</pre>
LLVM	<pre>{i8*, i8*, ; meminfo, parent i64, i64, ; nitems, itemsize i32*, ; data [1 x i64], [1 x i64]} ; shape, strides</pre>

BUILT-IN MODELS

- Data models can be re-used, subclassed by extensions
- A few examples built-in to Numba:

Model	Purpose	Examples
Primitive model	Primitive types	Integer values, floating point values, pointers
Opaque model	Objects passed as pointers	String literals, functions, exceptions,
Struct model	Objects represented as structures “under the hood”	Arrays, Unicode strings, Tuples
UniTuple model	Tuples of homogeneous type	UniTuple
Composite model	Objects with structure visible to the user	Records

- Use case in worked example

LOWERING

- Provide implementations of all supported operations
- Data flow handled by Numba:
 - Arguments passed in
 - Arguments returned
- 1 lowering per function and set of types
 - (Note: one lowering can handle multiple sets)
- Lowering function accepts:
 - **Context:** used for type-related operations
 - **Builder:** an llvmlite builder for building LLVM IR
 - **Signature** of the function - args and return type
 - **Arguments** to the function, LLVM IR values which will have been constructed earlier by other lowering functions
- Lowering function returns:
 - **LLVM IR** implementing the operation

```
@lower_builtin(<FUNCTION>, *argtypes)
def my_lowering_function(context, builder, signature, args):
    # Using context and builder
    # Generate llvm_ir implementing FUNCTION
    # for the provided signature
    # Operating on the given args
    return llvm_ir
```

LOWERING EXAMPLE - LEFT SHIFT ON INTS

- Compiling this function:

```
# x and y are int32 type
def f(x, y):
    return x << y
```

- Lowering for << on ints in Numba:

```
@lower_builtin(operator.lshift, Integer, Integer)
def int_shl_impl(context, builder, sig, args):
    [valty, amtty] = sig.args # int64, int64
    [val, amt] = args        # arg.x, arg.y
    val = context.cast(builder, val, valty, sig.return_type)
    amt = context.cast(builder, amt, amtty, sig.return_type)
    return builder.shl(val, amt)
```

- Generated LLVM IR:
(simplified)

```
define i32 @"f"(.., i32 %"arg.x", i32 %"arg.y")
{
    %.6" = sext i32 %"arg.y" to i64
    %.7" = sext i32 %"arg.x" to i64
    %.8" = shl i64 %.7", %.6"
}
```

BUILDING LLVM IR

- LlvmLite IR builder functions:

- Arithmetic / logical: add, sub, mul, sdiv, udiv, and, or, xor, not, shl, lshr, ashr, ctz, ctlz, ...
- Comparisons: Integer compare, floating point compare
- Branches and control flow: branch, if/then
- Function call / return: call, invoke, ret
- Aggregate operations: insert and extract values from structs and arrays
- Memory operations: load, store, atomics, ...
- Inline assembly: PTX, for CUDA

- (Many not listed)

- Full list at:

<https://llvmlite.readthedocs.io/en/latest/user-guide/ir/ir-builder.html>

WORKED EXAMPLE

Implementing a UDF compiler for Filigree

- Typing
- Data model
- Lowering
- Integration with Python API
- Full source, completed example, notes:
 - <https://github.com/gmarkall/numba-accelerated-udfs>

THE FILIGREE UDF COMPILER

- C++ headers / data structures / library + Python API
- Python API example that we're aiming to support:

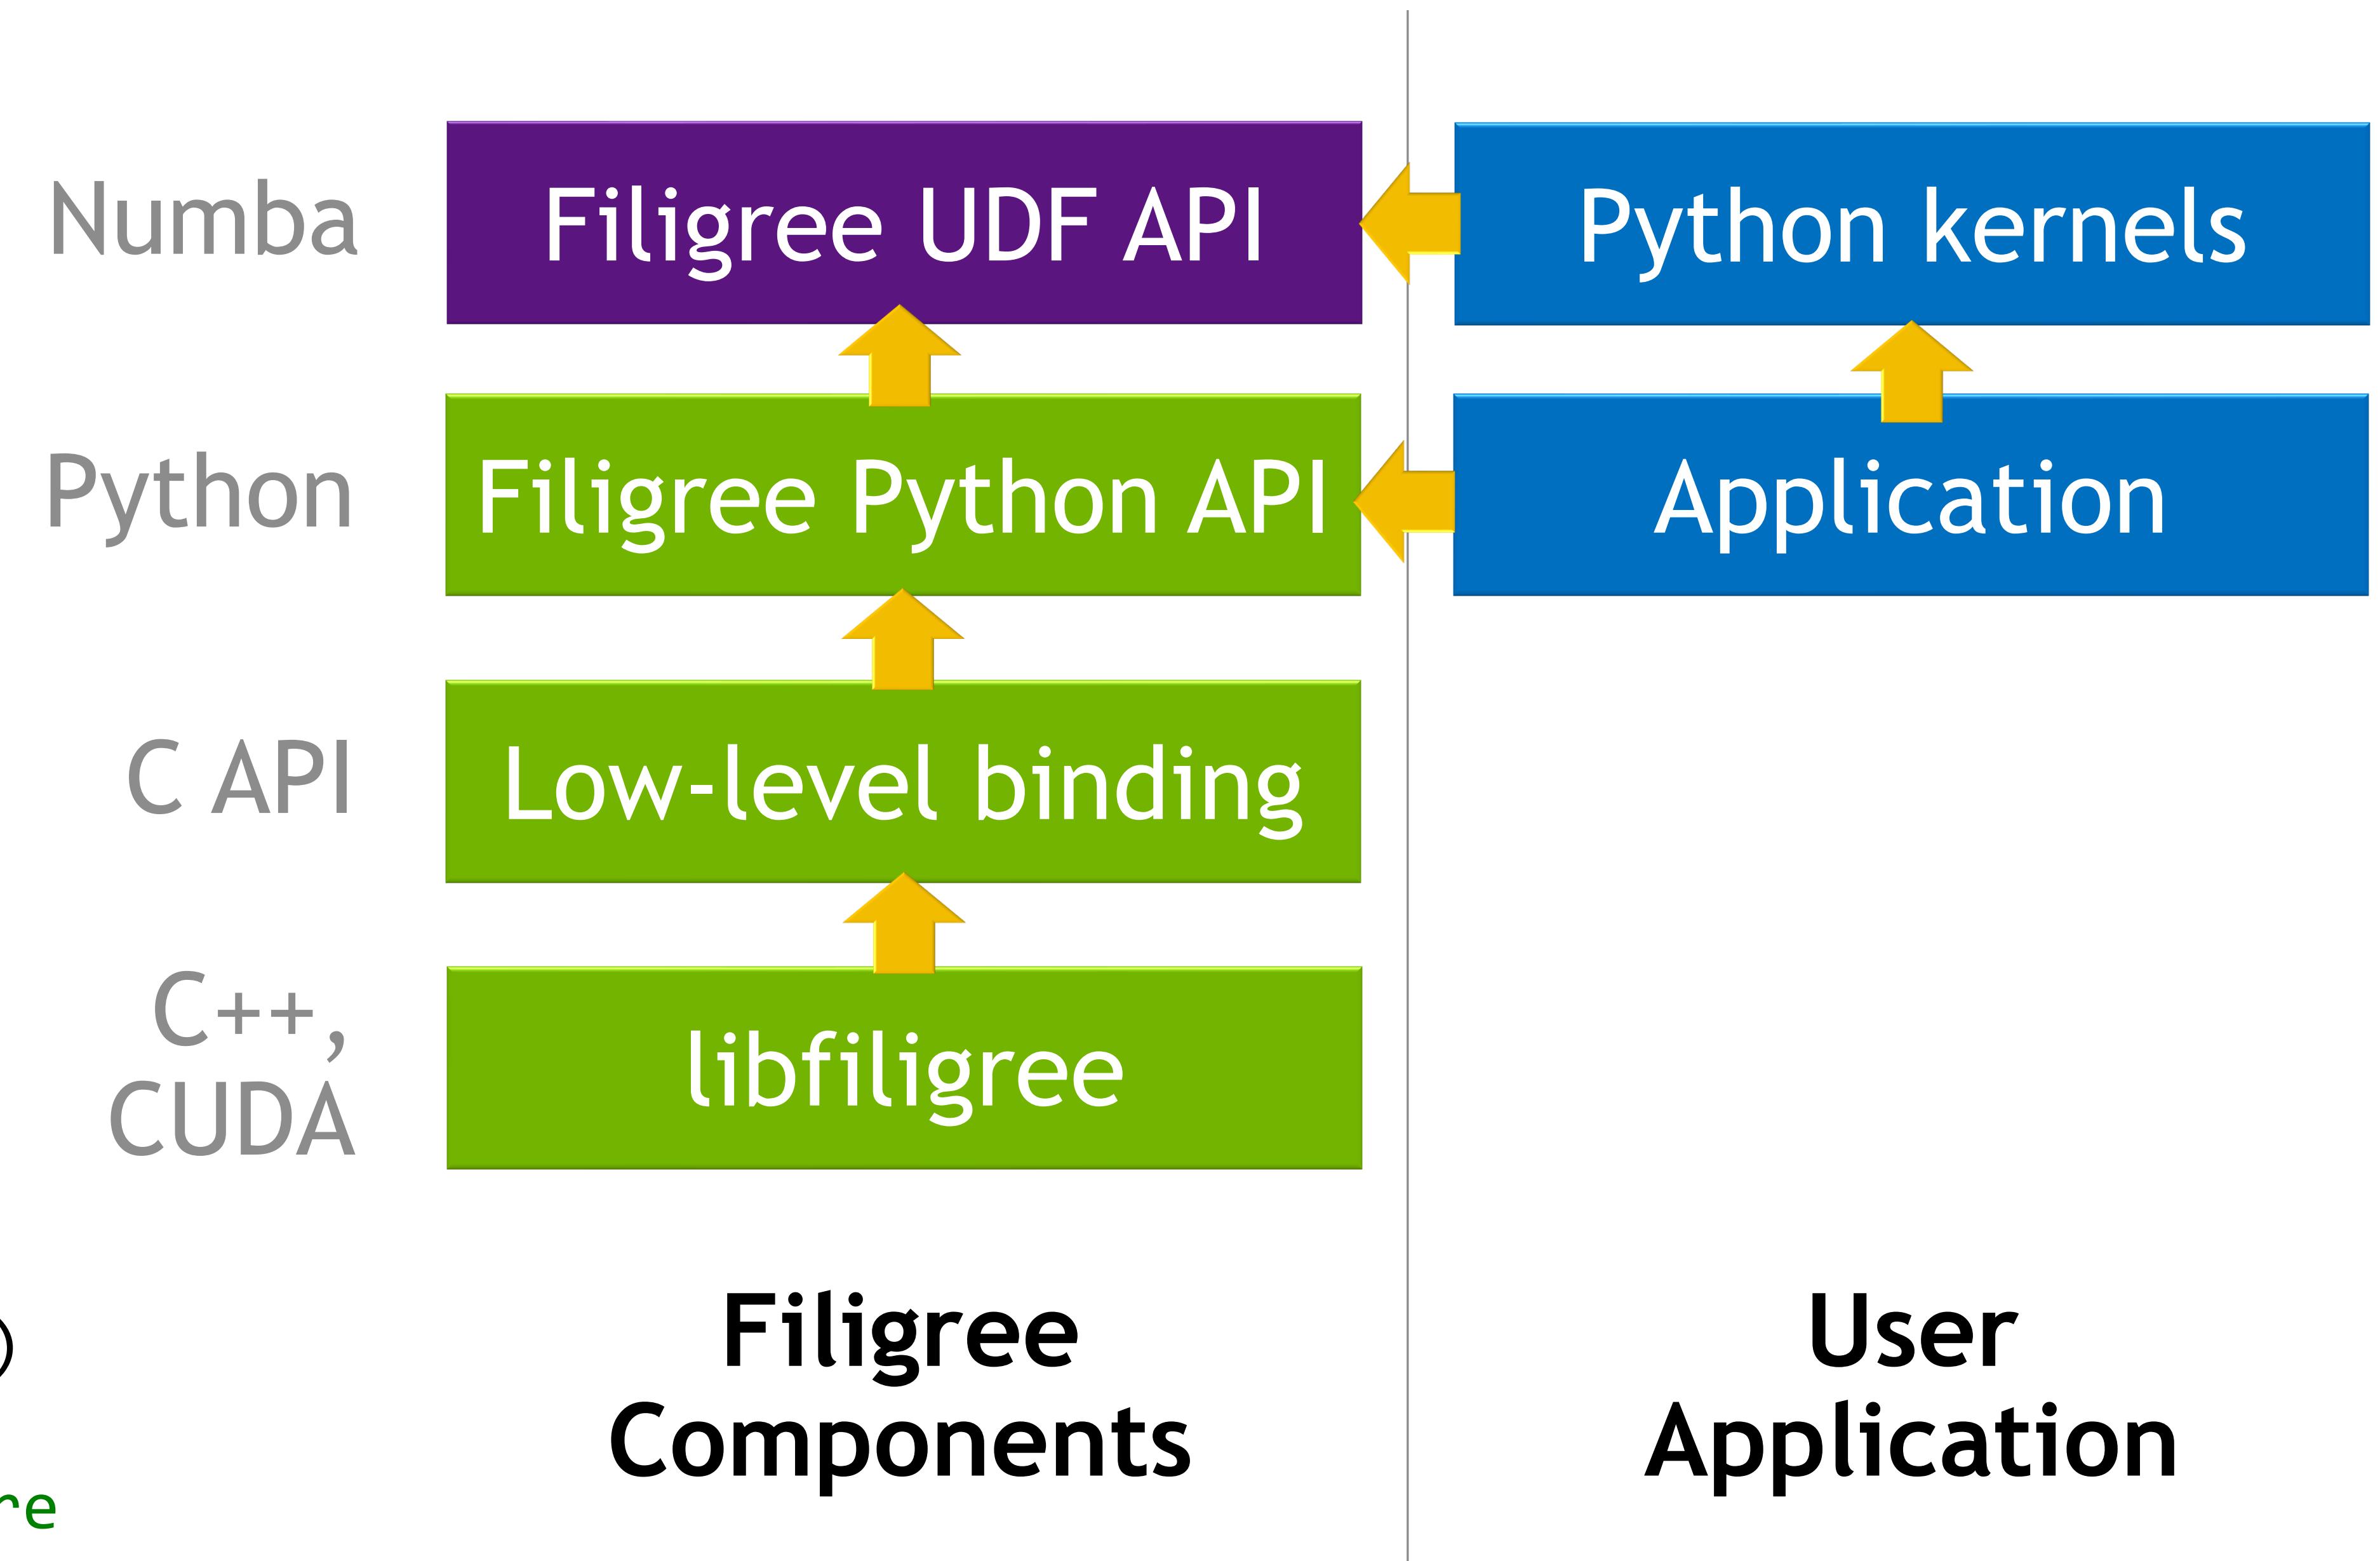
```
image = Image(path)

center_x = image.width / 2
center_y = image.height // 2
radius = max(center_x, center_y)

def highlight_center(pixel, x, y):
    distance_from_center = math.sqrt(
        abs(x - center_x)**2 + abs(y - center_y)**2)
    # Weight pixels according to distance from centre
    w = 1 - (distance_from_center / radius)

    # Apply weight to each RGB component
    return pixel.r * w, pixel.g * w, pixel.b * w, pixel.a

image.apply_located_pixel_udf(highlight_center)
image.to_pillow()
```



Filigree Components

TYPING: FILIGREE PIXEL STRUCTURES

```
// C++ Types (from ImageMagick): # Numba typing

typedef unsigned short Quantum;

struct PixelPacket {
    Quantum blue;
    Quantum green;
    Quantum red;
    Quantum opacity;
};

# Python PixelPacket
class PixelPacket:
    r: int
    g: int
    b: int
    a: int

    quantum = types.uint16

    class PixelPacket(types.Type):
        def __init__(self):
            super().__init__(name="PixelPacket")

        @cuda_registry.register_attr
        class PixelPacketAttrs(AttributeTemplate):
            key = pixel_packet

            def resolve_r(self, pp):
                return quantum

            def resolve_g(self, pp):
                return quantum

            def resolve_b(self, pp):
                return quantum

            def resolve_a(self, pp):
                return quantum
```

TYPING FILIGREE PIXEL STRUCTURES (POINTER)

```
// C++ Type (from Filigree)
class Image
{
public:
    Image(std::string filename, /* other params */);
    ~Image();

    void to_greyscale();
    void to_binary();
    void to_weighted_binary();

    // Other methods omitted...

private:
    rmm::mr::device_memory_resource *_mr;
    rmm::cuda_stream_view _stream;
    Magick::PixelPacket *_pixels;

    size_t _width;
    size_t _height;
    size_t _alloc_size;
};

# Numba type
class PixelPacketPointer(types.RawPointer):
    def __init__(self):
        super().__init__(name="PixelPacket")

    pixel_packet_pointer = PixelPacketPointer()

    # Reuse the r, g, b, a attributes
    @cuda_registry.register_attr
    class PixelPacketPointerAttrs(PixelPacketAttrs):
        key = pixel_packet_pointer

        # For dereferencing (pp a PixelPacketPointer):
        # pixel_packet = pp.values
        def resolve_values(self, pp):
            return pixel_packet

};
```

TYPING THE FILIGREE IMAGE CLASS

```
// C++ Type (from Filigree)
class Image
{
public:
    Image(std::string filename, /* ... */);
    ~Image();

    void to_greyscale();
    void to_binary();
    void to_weighted_binary();

    // Other methods omitted...

private:
    rmm::mr::device_memory_resource *_mr;
    rmm::cuda_stream_view _stream;

    Magick::PixelPacket *_pixels;

    size_t _width;
    size_t _height;
    size_t _alloc_size;
};

# Numba typing
class FiligreeImage(types.Type):
    def __init__(self):
        super().__init__(name="FiligreeImage")

    # E.g. for: pixel = image[x, y]
    @cuda_registry.register_global(operator.getitem)
    class Image_getitem(CallableTemplate):
        def generic(self):
            def typer(image, indices):
                if not isinstance(image, FiligreeImage):
                    return None
                if (not isinstance(indices, types.BaseTuple)
                    or len(indices) != 2):
                    return None
                return signature(pixel_packet_pointer, image, indices)

            return typer
```

DATA MODELS: PIXEL PACKET DATA MODEL

```
// C++ Types (from ImageMagick):                                # Numba data model

typedef unsigned short Quantum;

struct PixelPacket {
    Quantum blue;
    Quantum green;
    Quantum red;
    Quantum opacity;
};

@register_model(PixelPacket)
class PixelPacketModel(models.StructModel):
    def __init__(self, dmm, fe_type):
        members = [
            ("b", quantum),
            ("g", quantum),
            ("r", quantum),
            ("a", quantum)
        ]
        super().__init__(dmm, fe_type, members)
```

DATA MODELS: IMAGE DATA MODEL

```
// C++ Type (from Filigree)
class Image
{
public:
    Image(std::string filename, /* ... */);
    ~Image();

    void to_greyscale();
    void to_binary();
    void to_weighted_binary();

    // Other methods omitted...

private:
    rmm::mr::device_memory_resource *_mr;
    rmm::cuda_stream_view _stream;

    Magick::PixelPacket *_pixels;

    size_t _width;
    size_t _height;
    size_t _alloc_size;
};

# Numba data model
@register_model(FiligreeImage)
class FiligreeImageModel(models.StructModel):
    def __init__(self, dmm, fe_type):
        members = [
            ("pixel_ptr", pixel_packet_pointer),
            ("width", types.uint64),
            ("height", types.uint64),
        ]
        super().__init__(dmm, fe_type, members)
```

LOWERING PIXEL PACKET ATTRIBUTES

```
# Numba data model

@register_model(PixelPacket)
class PixelPacketModel(models.StructModel):
    def __init__(self, dmm, fe_type):
        members = [
            ("b", quantum),
            ("g", quantum),
            ("r", quantum),
            ("a", quantum)
        ]
        super().__init__(dmm, fe_type, members)

# Numba lowering

make_attribute_wrapper(PixelPacket, 'r', 'r')
make_attribute_wrapper(PixelPacket, 'g', 'g')
make_attribute_wrapper(PixelPacket, 'b', 'b')
make_attribute_wrapper(PixelPacket, 'a', 'a')

# User code:

r, g, b, a = pp.r, pp.g, pp.b, pp.a
```

LOWERING IMAGE GETITEM

- `cutils`:
 - Code generation utilities for common patterns.
- In this example:
 - `create_struct_proxy`
 - `unpack_tuple`
 - `gep_inbounds`

```
# User code
```

```
def fun(image):
    x, y = cuda.grid(2)
    pixel = image[x, y]
```

```
# Numba lowering

@cuda_lower(operatorgetitem, filigree_image, types.UniTupl)
def filigree_image_getitem(context, builder, sig, args):
    image_arg, index_arg = args
    image_ty, index_ty = sig.args
    image = cutils.create_struct_proxy(image_ty)(context, builder,
                                                value=image_arg)
    x, y = cutils.unpack_tuple(builder, index_arg, count=2)

    x = context.cast(builder, x, index_ty[0], types.uint64)
    y = context.cast(builder, y, index_ty[1], types.uint64)
    offset = builder.add(builder.mul(y, image.width), x)

    ptr = builder.bitcast(image.pixel_ptr,
                          context.get_value_type(pixel_packet).as_pointer())
    current_pixel = cutils.gep_inbounds(builder, ptr, offset)
    return current_pixel
```

LOWERING SETTING OF PIXEL ATTRIBUTES

```
# User code                                # Numba lowering

def set_blue(image):
    x, y = cuda.grid(2)
    pixel = image[x, y]
    pixel.b = 127

@cuda_lower_registry.lower_setattr_generic(pixel_packet_pointer)
def pixel_packet_pointer_set_attr(context, builder, sig, args, attr):
    # Load pixel data into struct proxy
    base_idx = context.get_constant(types.intp, 0)
    value_ptr = cgutils.gep_inbounds(builder, args[0], base_idx)
    data = builder.load(value_ptr)
    values = cgutils.create_struct_proxy(pixel_packet)(context, builder, value=data)

    # Set value. Alternative to:
    # if attr == 'r':
    #     values.r = args[1]
    # elif attr == 'g':
    #     values.g = args[1]
    setattr(values, attr, args[1])

    # Store struct value back to memory
    builder.store(values._getvalue(), value_ptr)
```

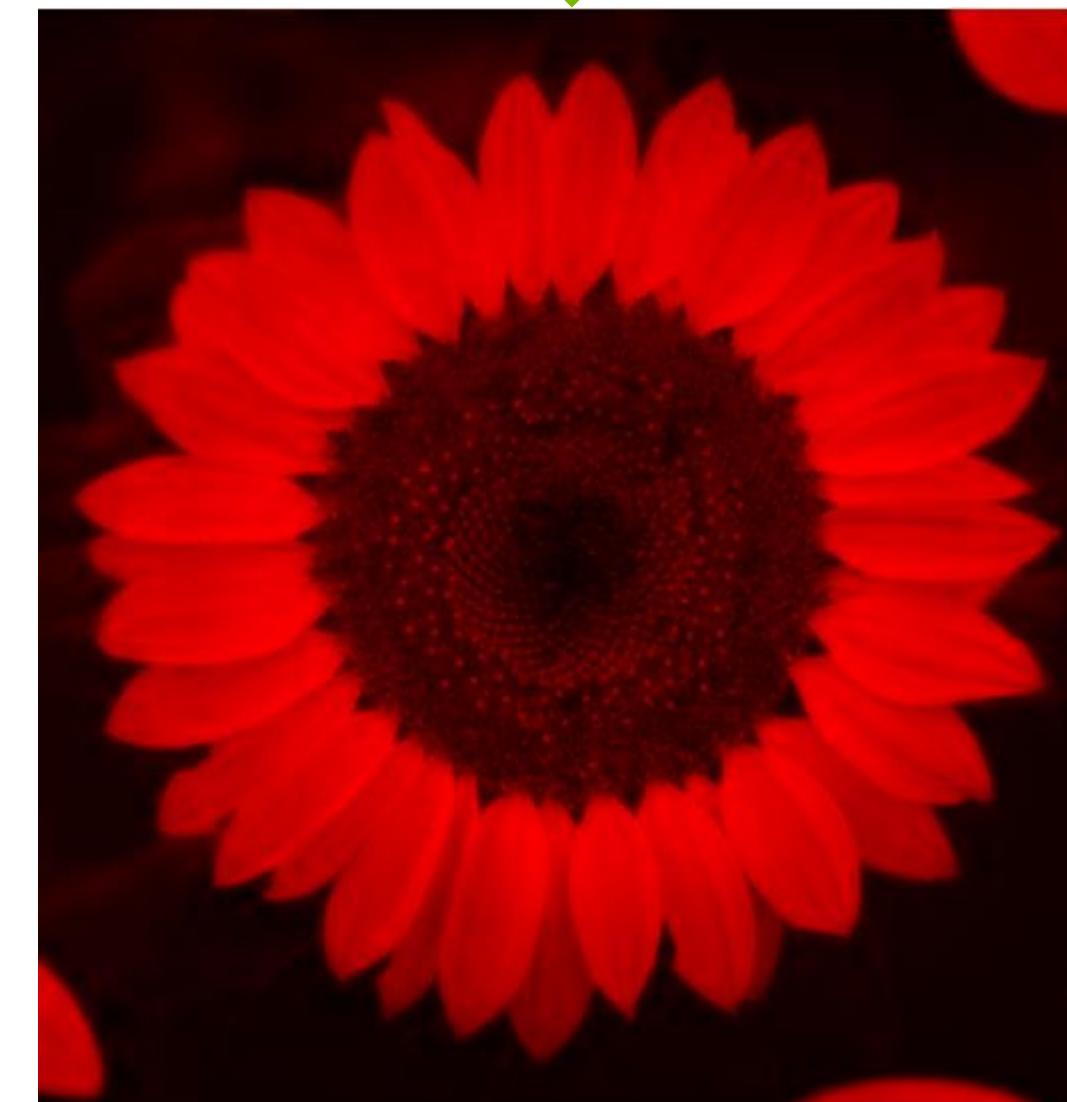
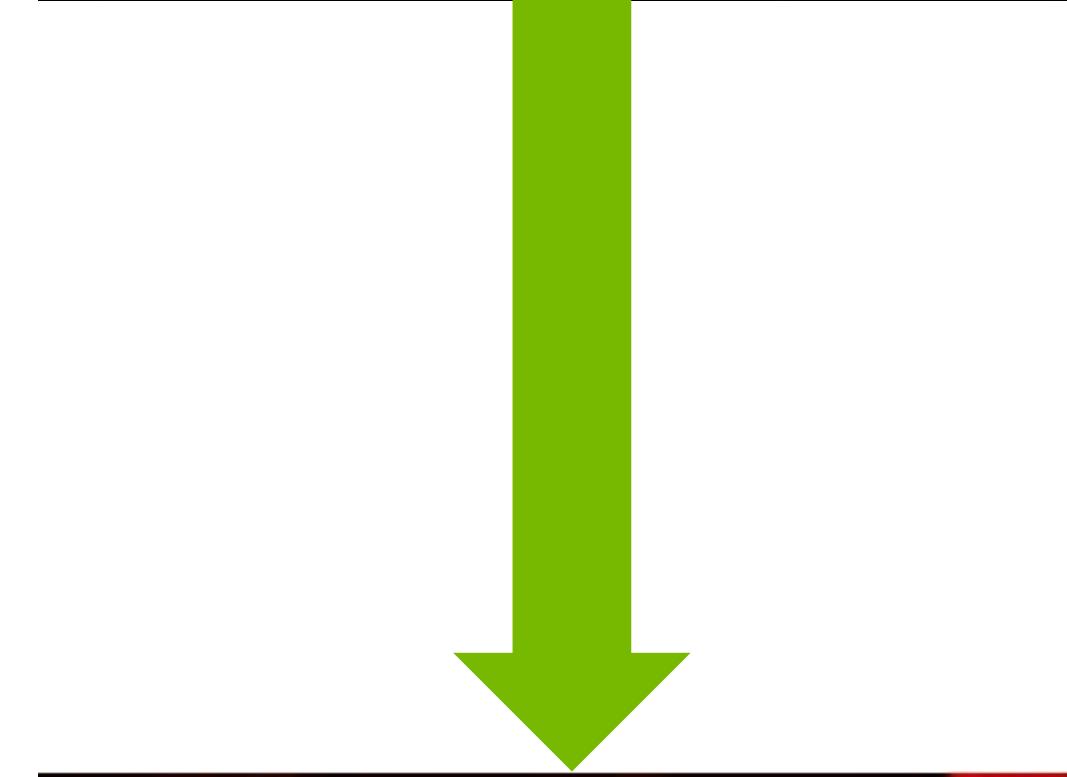
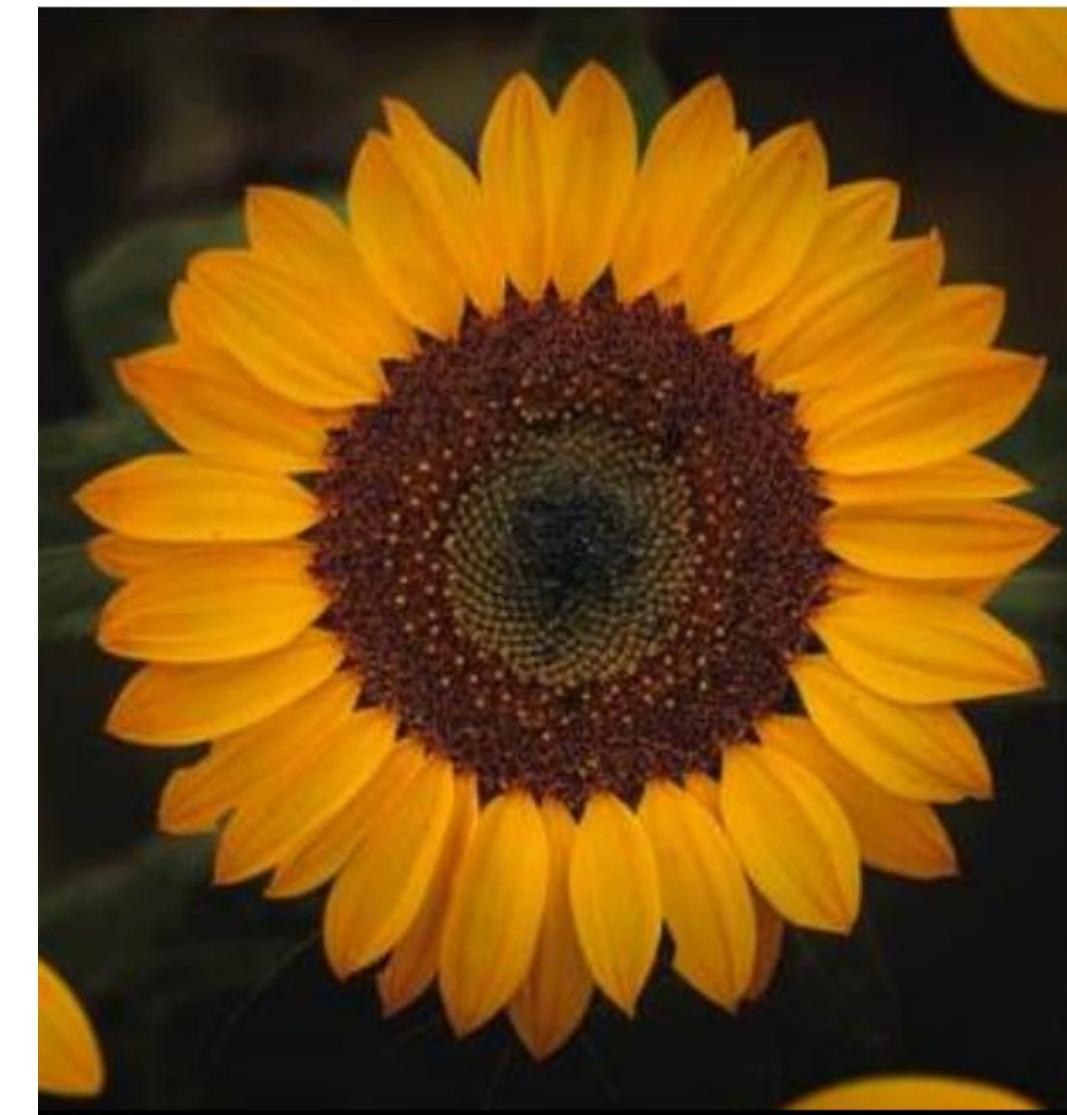
SUPPORTED USE CASE SO FAR

- Given our typing, lowering, and data model, we can compile:

```
@cuda.jit
def red_filter(pixel):
    # Strip out green and blue components
    return pixel.r, 0, 0, pixel.a
```

```
@cuda.jit
def my_transform(image):
    x, y = cuda.grid(2)
    pixel = image[x, y]
    r, g, b, a = red_filter(pixel.values)
    pixel.r = r
    pixel.g = g
    pixel.b = b
    pixel.a = a

image = Image('sunflower.png')
my_transform[grid_dim, block_dim](image)
```



- Need to add:

- Mechanism to pass the Image object to a kernel
- API for pixel-centric kernel launch

PYTHON LAUNCH API - PASSING IN IMAGES

```
# User code

def red_filter(pixel):
    # Strip out green and blue components
    return pixel.r, 0, 0, pixel.a

image = Image('sunflower.png')
Image.apply_located_pixel_udf(red_filter)

# Adaptor class

class FiligreeImageHandler:
    def prepare_args(self, ty, val, **kwargs):
        # Transform argument if it's an image
        # Otherwise, passthrough unchanged.
        if isinstance(val, api.Image):
            # Type + values match data model
            ty = types.UniTuple(types.uint64, 3)
            val = (val.pixel_ptr, val.width, val.height)
        return ty, val

filigree_image_handler = FiligreeImageHandler()

@cuda.jit(extensions=[filigree_image_handler])
```

```
# From Filigree Python API:

class Image:
    def __init__(self, filename):
        # _lib is the Python C API wrapper
        # for libfiligree
        self._lib_image = _lib.Image(filename)

    @property
    def pixel_ptr(self):
        return self._lib_image.get_pixel_ptr()

    @property
    def width(self):
        return self._lib_image.get_width()

    @property
    def height(self):
        return self._lib_image \
            .get_height()
```

PYTHON LAUNCH API - GENERATING THE APPLY KERNEL

```
# User code                                # API Implementation

def red_filter(pixel):
    # Strip out green and blue components
    return pixel.r, 0, 0, pixel.a

image = Image('sunflower.png')
Image.apply_located_pixel_udf(red_filter)

# API Implementation

def apply_located_pixel_udf(image, udf):
    device_function = cuda.jit(device=True)(udf)
    @cuda.jit(extensions=[filigree_image_handler])
    def apply_udf(image):
        x, y = cuda.grid(2)

        if x < image.width and y < image.height:
            pixel_ptr = image[x, y]
            r, g, b, a = device_function(pixel_ptr.values, x, y)
            pixel_ptr.r = r
            pixel_ptr.g = g
            pixel_ptr.b = b
            pixel_ptr.a = a

        nthreads_x = 16
        nthreads_y = 8
        nblocks_x = (image.width // nthreads_x) + 1
        nblocks_y = (image.width // nthreads_y) + 1
        grid_dim = (nblocks_x, nblocks_y)
        block_dim = (nthreads_x, nthreads_y)

        apply_udf[grid_dim, block_dim](image)
```

LAUNCH API DESIGN CONSIDERATIONS

- Caching options:
 - `functools.lru_cache`:
 - Simple, but fails to cache user-generated functions
 - Just add `@lru_cache` decorator to kernel generation
 - Bytecode-based caching:
 - More complex, handles users generating functions
 - Used in cuDF and Numba's on-disk cache
- Launch options:
 - Python: Used in Filigree, cuDF
 - C++: Used in PyOptiX, cuDF
 - Use `cuda.compile_ptx` to get PTX and invoke from C++
- See example repository for links

CONCLUDING POINTS

- Evaluation
- Summary
- Resources

PTX COMPARISON: RED FILTER KERNEL

```
.visible .entry my_transform(
    .param .u64 my_transform_param_0,
    .param .u64 my_transform_param_1,
    .param .u64 my_transform_param_2
)
{
    .reg .pred %p<4>;
    .reg .b16 %rs<2>;
    .reg .b32 %r<9>;
    .reg .b64 %rd<11>;

    ld.param.u64    %rd3, [my_transform_param_0];
    ld.param.u64    %rd4, [my_transform_param_1];
    ld.param.u64    %rd5, [my_transform_param_2];
    mov.u32    %r1, %tid.x;
    mov.u32    %r2, %ctaid.x;
    mov.u32    %r3, %ntid.x;
    mad.lo.s32 %r4, %r3, %r2, %r1;
    mov.u32    %r5, %ctaid.y;
    mov.u32    %r6, %ntid.y;
    mov.u32    %r7, %tid.y;
    mad.lo.s32 %r8, %r6, %r5, %r7;
    cvt.u64.u32    %rd1, %r4;
    setp.ge.u64    %p1, %rd1, %rd4;
    cvt.u64.u32    %rd2, %r8;
    setp.ge.u64    %p2, %rd2, %rd5;
    or.pred    %p3, %p1, %p2;
    @%p3 bra    $L__BB0_2;

    cvta.to.global.u64 %rd6, %rd3;
    mul.lo.s64    %rd7, %rd2, %rd4;
    add.s64     %rd8, %rd7, %rd1;
    shl.b64     %rd9, %rd8, 3;
    add.s64     %rd10, %rd6, %rd9;
    mov.u16     %rs1, 0;
    st.global.u16 [%rd10], %rs1;
    st.global.u16 [%rd10+2], %rs1;

$L__BB0_2:
    ret;
}

.visible .entry my_transform(
    .param .u64 my_transform_param_0,
    .param .u64 my_transform_param_1,
    .param .u64 my_transform_param_2
)
{
    .reg .pred %p<4>;
    .reg .b16 %rs<2>;
    .reg .b32 %r<9>;
    .reg .b64 %rd<11>;

    ld.param.u64    %rd3, [my_transform_param_0];
    ld.param.u64    %rd4, [my_transform_param_1];
    ld.param.u64    %rd5, [my_transform_param_2];
    mov.u32    %r1, %ntid.x;
    mov.u32    %r2, %ctaid.x;
    mov.u32    %r3, %tid.x;
    mad.lo.s32 %r4, %r2, %r1, %r3;
    mov.u32    %r5, %ntid.y;
    mov.u32    %r6, %ctaid.y;
    mov.u32    %r7, %tid.y;
    mad.lo.s32 %r8, %r6, %r5, %r7;
    cvt.u64.u32    %rd1, %r4;
    setp.gt.u64    %p1, %rd1, %rd4;
    cvt.u64.u32    %rd2, %r8;
    setp.gt.u64    %p2, %rd2, %rd5;
    or.pred    %p3, %p1, %p2;
    @%p3 bra    $L__BB0_2;

    mul.lo.s64    %rd6, %rd2, %rd4;
    add.s64     %rd7, %rd6, %rd1;
    cvta.to.global.u64 %rd8, %rd3;
    shl.b64     %rd9, %rd7, 3;
    add.s64     %rd10, %rd8, %rd9;
    mov.u16     %rs1, 0;
    st.global.u16 [%rd10], %rs1;
    st.global.u16 [%rd10+2], %rs1;

$L__BB0_2:
    ret;
}
```

PTX COMPARISON: HIGHLIGHT_CENTER DEVICE FUNCTION

```

.visible .func (.param .b32 func_retval0) highlight_center_py(
    .param .b64 highlight_center_py_param_0,
    .param .b32 highlight_center_py_param_1,
    .param .b32 highlight_center_py_param_2,
    .param .b32 highlight_center_py_param_3,
    .param .b32 highlight_center_py_param_4,
    .param .b32 highlight_center_py_param_5,
    .param .b32 highlight_center_py_param_6
)
{
    .reg .pred %p<3>;
    .reg .b16 %rs<5>;
    .reg .b32 %r<2>;
    .reg .f64 %fd<11>;
    .reg .b64 %rd<15>;

    ld.param.u64    %rd1, [highlight_center_py_param_0];
    ld.param.u16    %rs1, [highlight_center_py_param_1];
    ld.param.u16    %rs2, [highlight_center_py_param_2];
    ld.param.u16    %rs3, [highlight_center_py_param_3];
    ld.param.u16    %rs4, [highlight_center_py_param_4];
    ld.param.u32    %rd2, [highlight_center_py_param_5];
    ld.param.u32    %rd3, [highlight_center_py_param_6];
    add.s64        %rd4, %rd2, -155;
    setp.lt.u64    %p1, %rd2, 155;
    mov.u64        %rd5, 155;
    sub.s64        %rd6, %rd5, %rd2;
    selp.b64       %rd7, %rd6, %rd4, %p1;
    add.s64        %rd8, %rd3, -162;
    setp.lt.u64    %p2, %rd3, 162;
    mov.u64        %rd9, 162;
    sub.s64        %rd10, %rd9, %rd3;
    selp.b64       %rd11, %rd10, %rd8, %p2;
    mul.lo.s64    %rd12, %rd7, %rd7;
    mul.lo.s64    %rd13, %rd11, %rd11;
    add.s64        %rd14, %rd13, %rd12;
    cvt.rn.f64.s64 %fd1, %rd14;
    sqrt.rn.f64    %fd2, %fd1;
    div.rn.f64    %fd3, %fd2, 0dC064400000000000000;
    add.f64        %fd4, %fd3, 0d3FF000000000000;
    cvt.rn.f64.u16 %fd5, %rs3;
    mul.f64        %fd6, %fd4, %fd5;
    cvt.rn.f64.u16 %fd7, %rs2;
    mul.f64        %fd8, %fd4, %fd7;
    cvt.rn.f64.u16 %fd9, %rs1;
    mul.f64        %fd10, %fd4, %fd9;
    st.f64        [%rd1], %fd6;
    st.f64        [%rd1+8], %fd8;
    st.f64        [%rd1+16], %fd10;
    st.u16        [%rd1+24], %rs4;
    mov.u32        %r1, 0;
    st.param.b32   [func_retval0+0], %r1;
    ret;
}

.visible .func highlight_center(
    .param .b64 highlight_center_param_0,
    .param .align 2 .b8 highlight_center_param_1[8],
    .param .b32 highlight_center_param_2,
    .param .b32 highlight_center_param_3
)
{
    .reg .pred %p<3>;
    .reg .b16 %rs<5>;
    .reg .b32 %r<16>;
    .reg .f64 %fd<11>;
    .reg .b64 %rd<2>;

    ld.param.u32    %r1, [highlight_center_param_2];
    ld.param.u32    %r2, [highlight_center_param_3];
    ld.param.u16    %rs1, [highlight_center_param_1+6];
    ld.param.u16    %rs2, [highlight_center_param_1];
    ld.param.u16    %rs3, [highlight_center_param_1+2];
    ld.param.u16    %rs4, [highlight_center_param_1+4];

    add.s32        %r3, %r1, -155;
    mov.u32        %r4, 155;
    sub.s32        %r5, %r4, %r1;
    setp.lt.s32    %p1, %r3, 0;
    selp.b32       %r6, %r5, %r3, %p1;
    add.s32        %r7, %r2, -162;
    mov.u32        %r8, 162;
    sub.s32        %r9, %r8, %r2;
    setp.lt.s32    %p2, %r7, 0;
    selp.b32       %r10, %r9, %r7, %p2;
    mul.lo.s32    %r11, %r6, %r6;
    mad.lo.s32    %r12, %r10, %r10, %r11;
    cvt.rn.f64.u32 %fd1, %r12;
    sqrt.rn.f64    %fd2, %fd1;
    div.rn.f64    %fd3, %fd2, 0dC064400000000000000;
    add.f64        %fd4, %fd3, 0d3FF000000000000;
    cvt.rn.f64.u16 %fd5, %rs4;
    mul.f64        %fd6, %fd4, %fd5;
    cvt.rzi.u32.f64 %r13, %fd6;
    ld.param.u64    %rd1, [highlight_center_param_0];
    st.u16        [%rd1+4], %r13;
    cvt.rn.f64.u16 %fd7, %rs3;
    mul.f64        %fd8, %fd4, %fd7;
    cvt.rzi.u32.f64 %r14, %fd8;
    st.u16        [%rd1+2], %r14;
    cvt.rn.f64.u16 %fd9, %rs2;
    mul.f64        %fd10, %fd4, %fd9;
    cvt.rzi.u32.f64 %r15, %fd10;
    st.u16        [%rd1], %r15;
    st.u16        [%rd1+6], %rs1;
    ret;
}

```

SUMMARY AND RESOURCES

- Using Numba to compile Python code for CUDA solves the “*Impedance mismatch*” problem:
 - Python Users can *extend accelerated applications*
 - By writing their UDFs in Python they sustain their own *productivity* and *workflow*
 - And they can get *high performance code* generated equivalent to that from CUDA C++ with NVCC.
- Github Repository for this talk: <https://github.com/gmarkall/numba-accelerated-udfs>
- Other example extensions: <https://github.com/gmarkall/extending-numba-cuda>
 - [Jupyter Notebook](#) / [Quaternion Example](#) / [Interval Example](#)
- Application use cases:
 - cuDF Extension code: <https://github.com/rapidsai/cudf/tree/branch-22.04/python/cudf/cudf/core/udf>
 - PyOptiX Extension code: <https://github.com/gmarkall/PyOptiX/tree/gtc2022>
- The Life of a Numba Kernel: <https://github.com/gmarkall/life-of-a-numba-kernel/>
 - [Blog post](#) / [Jupyter Notebook](#)
- NVIDIA Numba CUDA tutorial:
 - Github repository: <https://github.com/numba/nvidia-cuda-tutorial>
 - All slides: <https://raw.githubusercontent.com/numba/nvidia-cuda-tutorial/main/numba-for-cuda-programmers-complete.pdf>
- Numba documentation:
 - Low-level extension API: <https://numba.readthedocs.io/en/stable/extending/low-level.html>
 - Notes on Numba’s architecture: <https://numba.readthedocs.io/en/stable/developer/repo-map.html>
- Contact:
 - Numba real-time chat: <https://gitter.im/numba/numba>
 - Numba Discourse forums: <https://numba.discourse.group/>
 - Email / Twitter: gmarkall@nvidia.com / <https://twitter.com/gmarkall>



SUMMARY AND RESOURCES

- Using Numba to compile Python code for CUDA solves the “*Impedance mismatch*” problem:
 - Python Users can *extend accelerated applications*
 - By writing their UDFs in Python they sustain their own *productivity* and *workflow*
 - And they can get *high performance code* generated equivalent to that from CUDA C++ with NVCC.
- Github Repository for this talk: <https://github.com/gmarkall/numba-accelerated-udfs>
- Other example extensions: <https://github.com/gmarkall/extending-numba-cuda>
 - [Jupyter Notebook](#) / [Quaternion Example](#) / [Interval Example](#)
- Application use cases:
 - cuDF Extension code: <https://github.com/rapidsai/cudf/tree/branch-22.04/python/cudf/cudf/core/udf>
 - PyOptiX Extension code: <https://github.com/gmarkall/PyOptiX/tree/gtc2022>
- The Life of a Numba Kernel: <https://github.com/gmarkall/life-of-a-numba-kernel/>
 - [Blog post](#) / [Jupyter Notebook](#)
- NVIDIA Numba CUDA tutorial:
 - Github repository: <https://github.com/numba/nvidia-cuda-tutorial>
 - All slides: <https://raw.githubusercontent.com/numba/nvidia-cuda-tutorial/main/numba-for-cuda-programmers-complete.pdf>
- Numba documentation:
 - Low-level extension API: <https://numba.readthedocs.io/en/stable/extending/low-level.html>
 - Notes on Numba’s architecture: <https://numba.readthedocs.io/en/stable/developer/repo-map.html>
- Contact:
 - Numba real-time chat: <https://gitter.im/numba/numba>
 - Numba Discourse forums: <https://numba.discourse.group/>
 - Email / Twitter: gmarkall@nvidia.com / <https://twitter.com/gmarkall>



SUMMARY AND RESOURCES

- Using Numba to compile Python code for CUDA solves the “*Impedance mismatch*” problem:
 - Python Users can *extend accelerated applications*
 - By writing their UDFs in Python they sustain their own *productivity* and *workflow*
 - And they can get *high performance code* generated equivalent to that from CUDA C++ with NVCC.
- Github Repository for this talk: <https://github.com/gmarkall/numba-accelerated-udfs>
- Other example extensions: <https://github.com/gmarkall/extending-numba-cuda>
 - [Jupyter Notebook](#) / [Quaternion Example](#) / [Interval Example](#)
- Application use cases:
 - cuDF Extension code: <https://github.com/rapidsai/cudf/tree/branch-22.04/python/cudf/cudf/core/udf>
 - PyOptiX Extension code: <https://github.com/gmarkall/PyOptiX/tree/gtc2022>
- The Life of a Numba Kernel: <https://github.com/gmarkall/life-of-a-numba-kernel/>
 - [Blog post](#) / [Jupyter Notebook](#)
- NVIDIA Numba CUDA tutorial:
 - Github repository: <https://github.com/numba/nvidia-cuda-tutorial>
 - All slides: <https://raw.githubusercontent.com/numba/nvidia-cuda-tutorial/main/numba-for-cuda-programmers-complete.pdf>
- Numba documentation:
 - Low-level extension API: <https://numba.readthedocs.io/en/stable/extending/low-level.html>
 - Notes on Numba’s architecture: <https://numba.readthedocs.io/en/stable/developer/repo-map.html>
- Contact:
 - Numba real-time chat: <https://gitter.im/numba/numba>
 - Numba Discourse forums: <https://numba.discourse.group/>
 - Email / Twitter: gmarkall@nvidia.com / <https://twitter.com/gmarkall>





NVIDIA®